

# DS 5110 – Lecture 2

## NumPy

Roi Yehoshua

# Agenda

---

- ▶ NumPy arrays
- ▶ Data types
- ▶ Indexing and slicing
- ▶ Universal functions
- ▶ Broadcasting
- ▶ Linear algebra

# NumPy

- ▶ NumPy (Numerical Python) is the fundamental package for numerical and statistical computing in Python
- ▶ Offers a high-performance, richly functional  $n$ -dimensional array type called **ndarray**
- ▶ Supports **vectorization**: a single operation can be carried out on an entire array
- ▶ NumPy is written in C to ensure that it runs as fast as possible
- ▶ Contains useful linear algebra and statistical analysis methods
- ▶ Many Python libraries depend on NumPy



# Creating Arrays From Existing Data

- ▶ NumPy is typically imported as **np** so that you can access its members with **np**.

```
import numpy as np
```

- ▶ NumPy provides various functions for creating arrays
- ▶ You can create an array from an existing list by passing it to the function **np.array()**:

```
a = np.array([2, 3, 5, 7, 11])
```

```
print(a)
```

```
[ 2  3  5  7 11]
```

- ▶ All the elements in an array must be of the same data type

# Creating Two-Dimensional Arrays

- ▶ To create a 2D array, you can pass a list of lists:

```
A = np.array([[1, 2, 3], [4, 5, 6]])
```

```
print(A)
```

```
[[1 2 3]
 [4 5 6]]
```

# Array Attributes

- ▶ The array object provides some useful attributes:
  - ▶ **dtype** – the type of the elements in the array
  - ▶ **ndim** – the array's number of dimensions
  - ▶ **shape** – a tuple specifying the size of the array along each dimension

```
A = np.array([[1, 2, 3], [4, 5, 6]])
```

```
A.dtype
```

```
dtype('int32')
```

```
A.ndim
```

```
2
```

```
A.shape
```

```
(2, 3)
```

# Array Element Type

- ▶ The data type of the array is deduced from the type of its elements:

```
a = np.array([1.0, 2.0, 3.0])  
a.dtype
```

```
dtype('float64')
```

- ▶ You can also explicitly set the data type using the optional **dtype** argument:

```
a = np.array([1, 2, 3], dtype='uint8')  
a.dtype
```

```
dtype('uint8')
```

# Common NumPy Data Types

Data Type	Description
int8	Integer in a single byte: -128 to 127
int16	Integer in 2 bytes: -32768 to 32767
int32	Integer in 4 bytes: -2147483648 to 2147483647
int64	Integer in 8 bytes: $-2^{63}$ to $2^{63} - 1$
uint8	Unsigned integer in a single byte: 0 to 255
uint16	Unsigned integer in 2 bytes: 0 to 65535
uint32	Unsigned integer in 4 bytes: 0 to 4294967295
uint64	Unsigned integer in 8 bytes: 0 to $2^{64} - 1$
float32	Single-precision, signed float: $\sim 10^{-38}$ to $\sim 10^{38}$ with $\sim 7$ decimal digits of precision
float64	Double-precision, signed float: $\sim 10^{-308}$ to $\sim 10^{308}$ with $\sim 15$ decimal digits of precision



# Filling Arrays with Specific Values

- ▶ The functions **zeros**, **ones** and **full** create arrays containing 0s, 1s, or a specified value
  - ▶ Their first argument is a tuple of integers specified the desired shape

```
a = np.zeros((3, 3)) # a matrix 3x3 of all zeros  
a
```

```
array([[0., 0., 0.],  
       [0., 0., 0.],  
       [0., 0., 0.]])
```

```
b = np.ones(5) # a 1-D array of all ones  
b
```

```
array([1., 1., 1., 1., 1.])
```

```
c = np.full((2, 2), 7) # a constant array 2x2 with values = 7  
c
```

```
array([[7, 7],  
       [7, 7]])
```

# Creating Arrays with Random Values

- ▶ The module `np.random` contains functions to generate random arrays
- ▶ **`np.random.random(shape)`** creates an array with numbers sampled randomly from the uniform distribution over  $[0, 1)$

```
np.random.random((3, 3))
```

```
array([[0.53938706, 0.48493127, 0.6644849 ],  
       [0.9244017 , 0.12693321, 0.0937883 ],  
       [0.11086629, 0.46859553, 0.77811775]])
```

- ▶ **`np.random.randint(low, [high], size)`** creates an array with integer numbers sampled from the interval  $[low, high)$ 
  - ▶ If *high* is not specified, then the sampled interval is  $[0, low)$

```
np.random.randint(1, 10, 5)
```

```
array([3, 1, 7, 2, 3])
```

# Creating Arrays with Random Values


- ▶ **`np.random.normal(loc=0.0, scale=1.0, size)`** samples numbers from the normal distribution with mean = *loc* and standard deviation = *scale*

```
np.random.normal(size=(3, 3))
```

```
array([[ -0.25329196, -0.0870325 ,  1.77494967],  
       [  1.36233034, -1.04551834, -0.16560166],  
       [  1.04324528, -2.18187512, -1.33373651]])
```

```
np.random.normal(100, 10, size=(3, 3))
```

```
array([[ 71.33579679,  88.46795444, 105.73253249],  
       [ 94.35509741, 102.48275919,  86.61720805],  
       [ 90.20942587, 103.56148269, 104.07765674]])
```



# Creating Arrays from Ranges

- ▶ **np.arange([start], stop, [step])** creates an array from a sequence of numbers
- ▶ Its arguments are the same as in Python's range() function
  - ▶ *start* – start of the sequence (default = 0)
  - ▶ *stop* – end of the sequence (not included in the sequence)
  - ▶ *step* – space between values (default = 1)
- ▶ In contrast to range() can also generate sequences of floating-point numbers

```
np.arange(5)
```

```
array([0, 1, 2, 3, 4])
```

```
np.arange(0.0, 1.0, 0.1)
```

```
array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9])
```

# Creating Arrays from Ranges

- ▶ **`np.linspace(start, stop, num)`** creates an array of evenly spaced numbers
  - ▶ *start* – start of the sequence
  - ▶ *stop* – end of the sequence (included in the sequence)
  - ▶ *num* – number of evenly spaced values to generate (default = 50)
- ▶ For example to get 10 equals spaced numbers between 0 and  $\pi$ :

```
np.linspace(0, np.pi, 10)
```

```
array([0.          , 0.34906585, 0.6981317 , 1.04719755, 1.3962634 ,  
       1.74532925, 2.0943951 , 2.44346095, 2.7925268 , 3.14159265])
```

# Reshaping an Array

- ▶ You can use the **reshape()** method to change the dimensions of your array
- ▶ The new shape should have the same number of elements as the original shape

```
np.arange(1, 13).reshape(3, 4)
```

```
array([[ 1,  2,  3,  4],  
       [ 5,  6,  7,  8],  
       [ 9, 10, 11, 12]])
```

- ▶ One of the dimensions may be specified as -1
  - ▶ Its size is inferred from the length of the array and the remaining dimensions

```
np.arange(1, 13).reshape(3, -1)
```

```
array([[ 1,  2,  3,  4],  
       [ 5,  6,  7,  8],  
       [ 9, 10, 11, 12]])
```

# Flattening an Array

- ▶ **ravel()** can be used to flatten a multidimensional array into one dimension:

```
a = np.arange(1, 10).reshape(3, 3)  
a
```

```
array([[1, 2, 3],  
       [4, 5, 6],  
       [7, 8, 9]])
```

```
a.ravel()
```

```
array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

# Indexing and Selecting Elements

---

- ▶ In NumPy there are different ways to access or change values in arrays
  - ▶ Indexing
  - ▶ Slicing
  - ▶ Fancy indexing
  - ▶ Boolean indexing / masking
  - ▶ And combinations thereof



# Indexing

- ▶ An array is indexed by a tuple of integers, e.g., `a[i, j]`
  - ▶ This is different from Python lists where we used double square brackets `a[i][j]`

```
a = np.arange(1, 13).reshape(3, 4)
a
```

```
array([[ 1,  2,  3,  4],
       [ 5,  6,  7,  8],
       [ 9, 10, 11, 12]])
```

```
a[1, 2]  # row 1, column 2
```

```
7
```

```
a[-1, -1]  # can use negative indexes
```

```
12
```

```
a[1]  # row 1
```

```
array([5, 6, 7, 8])
```

# Slicing

- ▶ Slicing allows you to select multiple sequential rows/columns
  - ▶ To select all the elements along one of the dimensions use colon :

```
a[0:2]  # the first two rows
```

```
array([[1, 2, 3, 4],  
       [5, 6, 7, 8]])
```

```
a[:, 0:2]  # the first two columns
```

```
array([[ 1,  2],  
       [ 5,  6],  
       [ 9, 10]])
```

```
a[0:2, 0:2]  # the first two rows and columns
```

```
array([[1, 2],  
       [5, 6]])
```

```
a[1, 1:]  # the second row and second column onwards
```

```
array([6, 7, 8])
```

# Class Exercise

- Use slicing to select the highlighted elements from the given matrix:

1	2	3
4	5	6
7	8	9
10	11	12

(a)

1	2	3
4	5	6
7	8	9
10	11	12

(b)

1	2	3
4	5	6
7	8	9
10	11	12

(c)

1	2	3
4	5	6
7	8	9
10	11	12

(d)

1	2	3
4	5	6
7	8	9
10	11	12

(e)

1	2	3
4	5	6
7	8	9
10	11	12

(f)

# Fancy Indexing

- ▶ You can select multiple non-sequential rows/columns by specifying a list of indexes

```
a = np.arange(1, 13).reshape(3, 4)  
a
```

```
array([[ 1,  2,  3,  4],  
       [ 5,  6,  7,  8],  
       [ 9, 10, 11, 12]])
```

```
a[[0, 2]]
```

```
array([[ 1,  2,  3,  4],  
       [ 9, 10, 11, 12]])
```

```
a[:, [1, 3]]
```

```
array([[ 2,  4],  
       [ 6,  8],  
       [10, 12]])
```

# Boolean Indexing (Masking)

- ▶ Boolean indexing is used to select elements of the array that satisfy some condition
- ▶ For example, to extract only the even numbers from a given array:

```
a = np.arange(5)
a[a % 2 == 0]
array([0, 2, 4])
```

- ▶ The mask creates an array of Boolean values:

```
a % 2 == 0
array([ True, False,  True, False,  True])
```

- ▶ The True elements indicate which elements from the array to return
- ▶ We can also use Boolean indexing directly:

```
a[[True, True, False, False, True]]
array([0, 1, 4])
```

# Changing Elements

- Changes to a sliced or indexed array are reflected in the original array, e.g.,

```
a = np.arange(10)
a
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
a[[2, 1, 8, 4]] = 99
a
```

```
array([ 0, 99, 99,  3, 99,  5,  6,  7, 99,  9])
```

```
a[a < 5] = -1
a
```

```
array([-1, 99, 99, -1, 99,  5,  6,  7, 99,  9])
```

# Array Operators

- ▶ NumPy provides many operators that perform operations on entire arrays
- ▶ These operators work **element-wise** (applied to every element in the array)
- ▶ For example, you can apply arithmetic operators between arrays and numeric values:

```
a = np.arange(1, 6)  
a
```

```
array([1, 2, 3, 4, 5])
```

```
a + 10
```

```
array([11, 12, 13, 14, 15])
```

```
a * 2
```

```
array([ 2,  4,  6,  8, 10])
```

```
a**3
```

```
array([ 1,  8, 27, 64, 125], dtype=int32)
```

# Array Operators

- ▶ You may also perform arithmetic operations between arrays of the same shape:

```
a = np.arange(1, 6)
```

```
a
```

```
array([1, 2, 3, 4, 5])
```

```
b = np.linspace(1.1, 5.5, 5)
```

```
b
```

```
array([1.1, 2.2, 3.3, 4.4, 5.5])
```

```
a * b
```

```
array([ 1.1,  4.4,  9.9, 17.6, 27.5])
```



# Slowness of Loops

- ▶ Vectorized operations execute significantly faster than corresponding list operations

```
def compute_reciprocals(values):  
    result = np.empty(len(values))  
    for i in range(len(values)):  
        result[i] = 1.0 / values[i]  
    return result  
  
big_array = np.random.randint(1, 100, size=1000000)  
%timeit compute_reciprocals(big_array)
```

1.53 s  $\pm$  25.2 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

```
%timeit (1.0 / big_array)
```

3.61 ms  $\pm$  16.6  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 100 loops each)

# Comparing Arrays

- ▶ You can compare arrays with scalars or with other arrays
- ▶ Comparisons are performed element-wise and produce an array of Boolean values

```
a = np.arange(1, 6)  
a
```

```
array([1, 2, 3, 4, 5])
```

```
a < 3
```

```
array([ True,  True, False, False, False])
```

```
2 * a == a**2
```

```
array([False,  True, False, False, False])
```

```
2 * a < a**2
```

```
array([False, False,  True,  True,  True])
```

# Comparing Arrays

- ▶ Comparing two floating-point arrays with `==` may give incorrect results
- ▶ Instead, use **`isclose(a, b)`** to check if the elements are “close” enough to each other

```
x = np.linspace(0, np.pi, 5)
```

```
np.sin(x)**2 == (1 - np.cos(x)**2)
```

```
array([ True, False,  True,  True, False])
```

```
np.isclose(np.sin(x)**2, 1 - np.cos(x)**2)
```

```
array([ True,  True,  True,  True,  True])
```

comparing with `==` gives  
incorrect results

# Logical Operations

- ▶ You can use the operators `&`, `|` and `~` to perform *and*, *or*, *not* between Boolean arrays

```
a = np.array([1, 2, 3, 4, 5])
```

```
(a > 1) & (a < 4)
```

```
array([False,  True,  True, False, False])
```

```
(a < 2) | (a > 4)
```

```
array([ True, False, False, False,  True])
```

```
~(a > 3)
```

```
array([ True,  True,  True, False, False])
```

Parentheses are required  
here because of operator  
precedence rules

# Universal Functions

- ▶ NumPy offers many **universal functions (ufuncs)** that perform various element-wise operations on arrays
- ▶ For example, we can calculate the square root of an array's values using **np.sqrt()**

```
a = np.array([1, 4, 9, 16, 25, 36])
```

```
np.sqrt(a)
```

```
array([1., 2., 3., 4., 5., 6.])
```

- ▶ Some of the ufuncs are called when you use operators like + or \*
  - ▶ e.g.,  $a + b$  is equivalent to **np.add(a, b)**

# Universal Functions

---

- ▶ Math
  - ▶ add, subtract, multiply, divide, remainder, exp, log, sqrt, power, and more
- ▶ Trigonometry
  - ▶ sin, cos, tan, hycot, arcsin, arccos, arctan, and more
- ▶ Bit manipulation
  - ▶ bitwise\_and, bitwise\_or, bitwise\_xor, invert, left\_shift, right\_shift
- ▶ Comparison
  - ▶ greater, greater\_equal, less, less\_equal, equal, not\_equal, logical\_and, logical\_or, and more
- ▶ Floating point
  - ▶ floor, ceil, isinf, isnan, fabs, trunc, and more
- ▶ For a full list see <https://numpy.org/doc/stable/reference/ufuncs.html>

# Example: Trigonometric Functions

- ▶ NumPy provides some useful trigonometric functions:

```
theta = np.linspace(0, np.pi, 3)  
theta
```

```
array([0.          , 1.57079633, 3.14159265])
```

```
np.sin(theta)
```

```
array([0.0000000e+00, 1.0000000e+00, 1.2246468e-16])
```

```
np.cos(theta)
```

```
array([ 1.000000e+00,  6.123234e-17, -1.000000e+00])
```

```
np.tan(theta)
```

```
array([ 0.0000000e+00,  1.6331239e+16, -1.2246468e-16])
```

# Example: Exponents and Logarithms

- ▶ Another common type of operation are exponentials and logarithms:

```
x = [1, 2, 3]  
np.exp(x)
```

```
array([ 2.71828183,  7.3890561 , 20.08553692])
```

```
x = [1, 2, 4, 10]  
np.log(x)  #  $\ln(x)$ 
```

```
array([0.          , 0.69314718, 1.38629436, 2.30258509])
```

```
np.log2(x)
```

```
array([0.          , 1.          , 2.          , 3.32192809])
```

```
np.log10(x)
```

```
array([0.          , 0.30103   , 0.60205999, 1.          ])
```



# Aggregations

- ▶ Aggregation functions allow you to get various statistics about your data
  - ▶ e.g., `sum()`, `min()`, `max()`, `mean()`, `var()`, `std()`
- ▶ By default these functions ignore the shape and use all elements in the calculations

```
grades = np.array([[87, 96, 70], [92, 87, 80],  
                  [84, 67, 90], [100, 81, 92]])
```

```
grades
```

```
array([[ 87,  96,  70],  
       [ 92,  87,  80],  
       [ 84,  67,  90],  
       [100,  81,  92]])
```

```
grades.max()
```

```
100
```


```
grades.mean()
```

```
85.5
```

# Calculations by Row or Column


- ▶ Many aggregation methods can be performed on specific array dimensions
- ▶ These methods receive an **axis** keyword argument that specifies which dimension to use in the calculation

1D array




shape: (4,)

2D array



shape: (2, 3)

3D array



shape: (4, 3, 2)

# Calculations by Row or Column

- ▶ For example, to calculate the average grade on each *exam* we can use `axis=0`:

```
grades.mean(axis=0)
```

```
array([90.75, 82.75, 83.  ])
```

- ▶ To calculate the average grade of each student we can use `axis=1`:

```
grades.mean(axis=1)
```

```
array([84.33333333, 86.33333333, 80.33333333, 91.  ])
```

# Missing Values

- ▶ **np.nan** (“not a number”) represents a missing value or the outcome of a calculation that is not well-defined (e.g., 0/0)

```
a = np.array([1, np.nan, 3, 4])  
a
```

```
array([ 1., nan,  3.,  4.])
```

- ▶ Do not test nans for equality (`np.nan == np.nan` is False), instead use **np.isnan()**

```
np.isnan(a)
```

```
array([False,  True, False, False])
```

- ▶ NaN is a bit like a virus – the result of any arithmetic with nan will also be a nan:

```
a + 1
```

```
array([ 2., nan,  4.,  5.])
```

```
a.sum()
```

```
nan
```



# Other Aggregation Functions

- ▶ NumPy provides many other aggregation functions
- ▶ Most aggregates have a NaN-safe version that ignores missing values

Function	NaN-Safe Version	Description
<code>np.sum()</code>	<code>np.nansum()</code>	Compute sum of elements
<code>np.prod()</code>	<code>np.nanprod()</code>	Compute product of elements
<code>np.mean()</code>	<code>np.nanmean()</code>	Compute mean of elements
<code>np.std()</code>	<code>np.nanstd()</code>	Compute standard variation
<code>np.var()</code>	<code>np.nanvar()</code>	Compute variance
<code>np.min()</code>	<code>np.nanmin()</code>	Find minimum value
<code>np.max()</code>	<code>np.nanxmax()</code>	Find maximum value
<code>np.argmin()</code>	<code>np.nanargmin()</code>	Find index of minimum value
<code>np.argmax()</code>	<code>np.nanargmax()</code>	Find index of maximum value
<code>np.cumsum()</code>	<code>np.nancumsum()</code>	Compute the cumulative sum of elements

# Argmin and Argmax

- ▶ Return the indices of the minimum/maximum values along an axis

```
a = np.random.randint(0, 10, (3, 3))  
a
```

```
array([[7, 6, 3],  
       [5, 2, 8],  
       [6, 9, 9]])
```

```
a.argmin()
```

```
4
```

```
a.argmin(axis=0)
```

```
array([1, 1, 0], dtype=int64)
```

```
a.argmin(axis=1)
```

```
array([2, 1, 0], dtype=int64)
```

# Counting Nonzero Entries

- ▶ **np.count\_nonzero()** can be used to count the True entries in a Boolean array
- ▶ The counting can also be done along rows or columns by using the **axis** argument

```
a = np.random.randint(100, size=(3, 4))  
a
```

```
array([[22, 72, 90, 92],  
       [ 0, 71, 77, 12],  
       [98, 25, 68, 13]])
```

```
np.count_nonzero(a % 2 == 0)
```

```
8
```

```
np.count_nonzero(a % 2 == 0, axis=0)
```

```
array([3, 1, 2, 2], dtype=int64)
```

```
np.count_nonzero(a % 2 == 0, axis=1)
```

```
array([4, 2, 2], dtype=int64)
```

# Truth Value Testing

- ▶ To quickly check whether any or all Boolean values are True, use **np.any()** or **np.all()**:

```
a = np.random.randint(100, size=(3, 4))  
a
```

```
array([[87, 99, 81, 78],  
       [84, 53, 81, 73],  
       [68, 51, 43, 47]])
```

```
# are there any values less than 10?  
np.any(a < 10)
```

False

```
# are all values greater than 20?  
np.all(a > 20)
```

True

```
# are all values in each row greater than 50?  
np.all(a > 50, axis=1)
```

```
array([ True,  True, False])
```



# Class Exercise

---

- ▶ Write a function that gets a 2D NumPy array and a number, and returns how many rows in the array contain the given number
- ▶ For example, given the array 

```
[[5 2 7 6]  
 [8 3 0 3]  
 [1 7 3 7]  
 [0 2 1 2]]
```

and the number 7, the function should return 2

# Sorting an Array

- ▶ **np.sort()** returns a sorted version of the array without modifying it:

```
a = np.array([3, 1, 4, 2, 5])  
np.sort(a)
```

```
array([1, 2, 3, 4, 5])
```

```
a
```

```
array([3, 1, 4, 2, 5])
```

- ▶ To sort the array in-place, you can instead use the **sort()** method of arrays:

```
a.sort()  
a
```

```
array([1, 2, 3, 4, 5])
```

# Sorting an Array

- ▶ You can sort along specific rows or columns of a 2D array using the `axis` argument

```
a = np.random.randint(10, size=(4, 6))  
a
```

```
array([[0, 3, 1, 5, 8, 3],  
       [9, 1, 5, 9, 1, 6],  
       [4, 5, 4, 0, 8, 9],  
       [9, 7, 5, 0, 5, 6]])
```

```
# Sort each column of a  
np.sort(a, axis=0)
```

```
array([[0, 1, 1, 0, 1, 3],  
       [4, 3, 4, 0, 5, 6],  
       [9, 5, 5, 5, 8, 6],  
       [9, 7, 5, 9, 8, 9]])
```

```
# Sort each row of a  
np.sort(a, axis=1)
```

```
array([[0, 1, 3, 3, 5, 8],  
       [1, 1, 5, 6, 9, 9],  
       [0, 4, 4, 5, 8, 9],  
       [0, 5, 5, 6, 7, 9]])
```

# Sorting an Array

- ▶ Although Python has built-in `sort()` and `sorted()` functions for lists, NumPy's **`np.sort()`** function turns out to be much more efficient:

```
big_array = np.random.randint(1, 100, size=1000000)
%timeit sorted(big_array)
```

308 ms  $\pm$  5.53 ms per loop (mean  $\pm$  std. dev. of 7 runs, 1 loop each)

```
%timeit np.sort(big_array)
```

30.7 ms  $\pm$  906  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 10 loops each)

- ▶ By default `np.sort()` uses **quicksort**, though mergesort and heapsort are also available

# Sorting an Array

- ▶ **np.argsort()** returns the *indexes* that would sort an array rather than the sorted elements themselves:

```
a = [7, 3, 10, 2, 8]  
np.argsort(a)
```

```
array([3, 1, 0, 4, 2], dtype=int64)
```

# Searching

- ▶ **np.where(*condition*)** returns the indices of all the elements that satisfy the condition:

```
a = np.array([7, 2, 5, 1, 4, 6, 3])  
np.where(a > 4)
```

```
(array([0, 2, 5], dtype=int64),)
```

- ▶ In case of a 2D array, it returns a tuple of the row and the column indices of the elements that satisfy the condition:

```
a = np.arange(9).reshape(3, 3)  
a
```

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

```
np.where(a > 4)
```

```
(array([1, 2, 2, 2], dtype=int64), array([2, 0, 1, 2], dtype=int64))
```

# Unique Values

- ▶ To find the unique values in an array, you can use **np.unique()**:

```
a = np.array([1, 2, 6, 4, 2, 3, 2, 6])  
np.unique(a)
```

```
array([1, 2, 3, 4, 6])
```

- ▶ If you set the argument *return\_counts* to True, the function will also return the number of times each unique item appears in the array:

```
values, counts = np.unique(a, return_counts=True)  
counts
```

```
array([1, 3, 1, 1, 2], dtype=int64)
```

# Adding Elements to an Array

- ▶ **`np.append(arr, values, axis=None)`** appends values to the end of an array
  - ▶ *values* must be of the correct shape (the same shape as *arr*, excluding *axis*)
  - ▶ *axis* specifies the axis along which to insert *values*

```
a = np.array([1, 2, 3])  
np.append(a, [4, 5])
```

```
array([1, 2, 3, 4, 5])
```

```
b = np.arange(9).reshape(3, 3)  
b
```

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

```
np.append(b, [[9, 10, 11]], axis=0)
```

```
array([[ 0,  1,  2],  
       [ 3,  4,  5],  
       [ 6,  7,  8],  
       [ 9, 10, 11]])
```



# Joining Arrays

- ▶ NumPy provides several functions to merge two arrays
- ▶ **np.vstack(*tuple of arrays*)** stacks the arrays vertically (row-wise)

```
a = np.array([[1, 2, 3], [4, 5, 6]])  
b = np.array([[7, 8, 9], [10, 11, 12]])  
  
np.vstack((a, b))
```

```
array([[ 1,  2,  3],  
       [ 4,  5,  6],  
       [ 7,  8,  9],  
       [10, 11, 12]])
```

- ▶ **np.hstack(*tuple of arrays*)** stacks the arrays horizontally (column-wise)

```
np.hstack((a, b))
```

```
array([[ 1,  2,  3,  7,  8,  9],  
       [ 4,  5,  6, 10, 11, 12]])
```

# Broadcasting

- ▶ Broadcasting allows operations to be performed on arrays of different shapes
- ▶ The smaller array is “broadcast” across the larger array to make them compatible
- ▶ For example, we can multiply a  $2 \times 3$  matrix by a one-dimensional array of 3 elements:

```
a = np.arange(1, 7).reshape(2, 3)
a
```

```
array([[1, 2, 3],
       [4, 5, 6]])
```

```
b = np.array([2, 4, 6])
```

```
a * b
```


```
array([[ 2,  8, 18],
       [ 8, 20, 36]])
```

- ▶ NumPy treats array b as if it were the matrix  $\begin{bmatrix} 2 & 4 & 6 \\ 2 & 4 & 6 \end{bmatrix}$

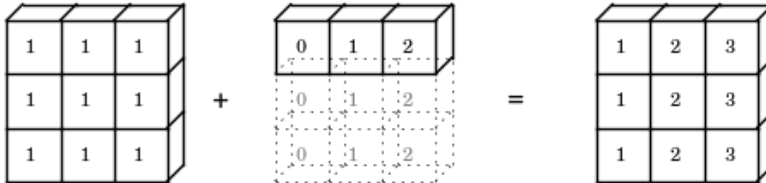
# Broadcasting Rules

- ▶ **Rule 1:** If the two arrays have different numbers of dimensions, the shape of the smaller array is **padded** with ones on its left side
- ▶ **Rule 2:** The arrays are **compatible** if all their dimensions are equal or one of them is 1
- ▶ **Rule 3:** If the arrays are compatible and their shapes don't match in one of the dimensions, the array with dimension of 1 is stretched to match the other array

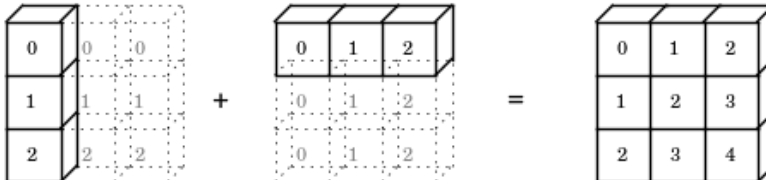
`np.arange(3) + 5`



`np.ones((3, 3)) + np.arange(3)`



`np.arange(3).reshape((3, 1)) + np.arange(3)`



# Broadcasting Rules

- ▶ When two arrays have shapes that don't support broadcasting, a `ValueError` occurs:

```
a = np.arange(1, 7).reshape(3, 2)
b = np.arange(3)

a + b
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-24-323d07b6b10f> in <module>
      2 b = np.arange(3)
      3
----> 4 a + b
```

**ValueError:** operands could not be broadcast together with shapes (3,2) (3,)

- ▶ The shapes of the arrays are: `a.shape = (3, 2)`, `b.shape = (3, )`
- ▶ By rule 1, we first pad the shape of `a` with ones: `a.shape -> (1, 3)`
- ▶ `a` and `b` are incompatible, since their second dimensions are different and none of them is 1

# Class Exercise

- Predict the result of the following operations:

```
x = np.arange(4)
xx = x.reshape(4, 1)
y = np.ones(5)
z = np.ones((3, 4))

print(xx + y)
print(x + z)
print(x + y)
```

# Matrices

- ▶ NumPy contains special methods to create matrices of specific types
- ▶ **np.eye( $N$ )** creates an identity matrix of size  $N \times N$

```
np.eye(3)
```

```
array([[1., 0., 0.],  
       [0., 1., 0.],  
       [0., 0., 1.]])
```

- ▶ **np.diag( $v$ )** creates a diagonal matrix from the 1-D array  $v$ :

```
np.diag([1, 2, 3, 4])
```

```
array([[1, 0, 0, 0],  
       [0, 2, 0, 0],  
       [0, 0, 3, 0],  
       [0, 0, 0, 4]])
```

# Transpose of a Matrix

- ▶ The transpose of a matrix results from “flipping” the rows and columns

$$(A^T)_{ij} = A_{ji}$$

- ▶ In NumPy, you can use the attribute **.T** to get the transposed matrix

```
A = np.arange(9).reshape(3, 3)  
A
```

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

```
A.T
```

```
array([[0, 3, 6],  
       [1, 4, 7],  
       [2, 5, 8]])
```

# Dot Product and Matrix Multiplication

- ▶ **`np.dot(a, b)`** computes the dot product of arrays  $a$  and  $b$ 
  - ▶ If both  $a$  and  $b$  are 1-D arrays, it is the **dot product** of vectors
  - ▶ If  $a$  or  $b$  is a 2-D array, it is **matrix multiplication**, but using  $a @ b$  is preferred

```
x = np.array([1, 2])
y = np.array([3, 4])

np.dot(x, y)  # 1 * 3 + 2 * 4
```

11

```
A = np.array([[4, 1], [2, 2]])
B = np.array([[1, 0], [0, 1]])

A @ B  # matrix multiplication
```

```
array([[4, 1],
       [2, 2]])
```

```
A @ x  # matrix-vector multiplication
```

```
array([6, 6])
```



# Linear Algebra

- ▶ The **numpy.linalg** module contains additional functions for working with matrices
- ▶ For example, **np.linalg.inv()** computes the inverse of a matrix:

```
a = np.array([[1, 2], [3, 4]])  
np.linalg.inv(a)
```

```
array([[ -2. ,  1. ],  
       [ 1.5, -0.5]])
```

- ▶ If the matrix is not invertible, then a `LinAlgError` exception is raised
- ▶ Full list of functions: <https://numpy.org/doc/stable/reference/routines.linalg.html>

# System of Linear Equations

- ▶ The set of linear equations:
$$\begin{aligned}a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n &= b_1 \\a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n &= b_2 \\&\vdots \\a_{n1}x_1 + a_{n2}x_2 + \dots + a_{nn}x_n &= b_n\end{aligned}$$

- ▶ Can be expressed as the matrix equation  $A\mathbf{x} = \mathbf{b}$ :

$$A\mathbf{x} = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix} = \mathbf{b}$$

- ▶ This system has a unique solution if only if  $A$  is non-singular
- ▶ In this case, we can write the solution as  $\mathbf{x} = A^{-1}\mathbf{b}$

# System of Linear Equations

- ▶ **`np.linalg.solve(A, b)`** computes the solution of the linear matrix equation  $Ax = b$ 
  - ▶ If no unique solution exists (for nonsquare or singular matrix  $A$ ), a `LinAlgError` is raised
- ▶ For example, let's find a solution to the following system of equations:

$$3x - 2y = 8$$

$$-2x + y - 3z = -20$$

$$4x + 6y + z = 7$$

```
A = np.array([[3, -2, 0],  
              [-2, 1, -3],  
              [4, 6, 1]])  
b = np.array([8, -20, 7])  
np.linalg.solve(A, b)
```

```
array([ 2., -1.,  5.])
```

```
np.linalg.inv(A) @ b
```

```
array([ 2., -1.,  5.])
```

# The Determinant

- ▶ The **determinant** of a square matrix  $A \in \mathbb{R}^{n \times n}$  is a function denoted by  $|A|$  or  $\det(A)$
- ▶ In the case of a  $2 \times 2$  matrix the determinant is computed by:

$$\begin{vmatrix} a & b \\ c & d \end{vmatrix} = ad - bc$$

- ▶ The general recursive formula for the determinant is:

$$\begin{aligned} |A| &= \sum_{i=1}^n (-1)^{i+j} a_{ij} |A_{\setminus i, \setminus j}| \quad (\text{for any } j \in 1, \dots, n) \\ &= \sum_{j=1}^n (-1)^{i+j} a_{ij} |A_{\setminus i, \setminus j}| \quad (\text{for any } i \in 1, \dots, n) \end{aligned}$$

- ▶  $A_{\setminus i, \setminus j}$  is the matrix that results from deleting the  $i$ th row and  $j$ th column from  $A$

# The Determinant

- ▶ For example, let's compute the determinant of the matrix  $A = \begin{pmatrix} 1 & 2 & 3 \\ 0 & 1 & 5 \\ 5 & 6 & 0 \end{pmatrix}$

$$|A| = 1 \cdot (1 \cdot 0 - 5 \cdot 6) - 2 \cdot (0 \cdot 0 - 5 \cdot 5) + 3 \cdot (0 \cdot 6 - 1 \cdot 5) = -30 + 50 - 15 = 5$$

- ▶ In Python, you can use the function **np.linalg.det()** to compute the determinant:

```
A = np.array([[1, 2, 3],  
              [0, 1, 5],  
              [5, 6, 0]])
```

```
np.linalg.det(A)
```


```
4.999999999999995
```

# Eigenvalues and Eigenvectors

- ▶ Given a square matrix  $A \in \mathbb{R}^{n \times n}$ , we say that  $\lambda \in \mathbb{C}$  is an **eigenvalue** of  $A$  and  $\mathbf{x} \in \mathbb{C}^n$  is its corresponding **eigenvector** if

$$A\mathbf{x} = \lambda\mathbf{x}, \quad \mathbf{x} \neq 0$$

- ▶ Intuitively, this definition means that multiplying  $A$  by the vector  $\mathbf{x}$  results in a new vector that points in the same direction as  $\mathbf{x}$ , but scaled by a factor  $\lambda$



- ▶ An eigenvector with its associated eigenvalue is called an **eigenpair**

# Eigenvalues and Eigenvectors

---

- ▶ To find the eigenpairs of a matrix  $A$ , we can rewrite the equation above as follows:

$$(A - \lambda I)\mathbf{x} = 0, \quad \mathbf{x} \neq 0$$

- ▶ But  $(A - \lambda I)\mathbf{x} = 0$  has non-zero solution to  $\mathbf{x}$  if and only if  $A - \lambda I$  is singular, i.e.,

$$|A - \lambda I| = 0$$

- ▶ We can use the definition of the determinant to expand this expression into a (very large) polynomial in  $\lambda$ , where  $\lambda$  will have degree  $n$
- ▶ This polynomial is often called the **characteristic polynomial** of the matrix  $A$

# Eigenvalues and Eigenvectors: Example

- ▶ **np.linalg.eig(A)** computes the eigenvalues and eigenvectors of the matrix  $A$ 
  - ▶ The eigenvectors are returned as normalized column vectors

```
A = np.array([[2, 1, 0],  
              [1, 2, 1],  
              [0, 1, 2]])
```

```
eigen_vals, eigen_vecs = np.linalg.eig(A)  
eigen_vals
```

```
array([3.41421356, 2.          , 0.58578644])
```

```
eigen_vecs
```

```
array([[ -5.00000000e-01,  7.07106781e-01,  5.00000000e-01],  
       [ -7.07106781e-01,  4.05925293e-16, -7.07106781e-01],  
       [ -5.00000000e-01, -7.07106781e-01,  5.00000000e-01]])
```