



Northeastern
University

DS 5110 – Lecture 6

SQL Part II

Roi Yehoshua



Agenda

- ▶ Subqueries
- ▶ Common table expressions (CTEs)
- ▶ Recursive queries
- ▶ Query execution plans
- ▶ Views
- ▶ Stored procedures
- ▶ Transactions
- ▶ Indexes
- ▶ Working with SQL from Python



Subqueries

- ▶ A subquery is a SELECT query nested inside another query
- ▶ Subqueries allow you to perform SQL operations in multiple steps
 - ▶ The subquery is executed before its parent query and its results are passed to the outer query
- ▶ A subquery may be nested inside
 - ▶ The SELECT clause
 - ▶ The FROM clause
 - ▶ The WHERE clause
 - ▶ The HAVING clause
- ▶ Subqueries must be enclosed inside brackets ()
- ▶ Subqueries can return an individual value or a list of records



Subqueries

- ▶ Example: Find the name of the instructor with the highest salary
- ▶ This query cannot be written without sorting the whole table
- ▶ The way to write it with a subquery:

A screenshot of MySQL Workbench showing a query editor and a results grid. The query editor contains the following SQL code:

```
1 •  SELECT name, salary
2   FROM instructor
3 WHERE salary = (
4     SELECT MAX(salary)
5     FROM instructor
6   );
7
```

The results grid shows one row of data:

	name	salary
▶	Einstein	95000.00



Correlated Subqueries

- ▶ A correlated subquery references one or more columns from the outer query
- ▶ A correlated subquery is executed once for each candidate row in the outer query
- ▶ Example: Find all instructors whose salary is higher than the average salary **in their own department**

The screenshot shows a MySQL Workbench interface. At the top is a toolbar with various icons. Below it is a SQL editor window containing the following code:

```
1 • SELECT i1.name, i1.salary
2   FROM instructor i1
3 WHERE salary > (
4     SELECT AVG(salary)
5     FROM instructor i2
6     WHERE i2.dept_name = i1.dept_name
7 );
```

Below the SQL editor is a "Result Grid" window displaying the query results:

	name	salary
▶	Wu	90000.00
	Einstein	95000.00
	Califieri	62000.00
	Brandt	92000.00



Multiple-Row Subqueries

- ▶ If the subquery returns more than one row, you cannot use it with the regular comparison operators like `<`, `>`, `=`
- ▶ There are 4 additional operators that allow to use the results of these subqueries:
 - ▶ `IN` / `NOT IN`
 - ▶ `ALL`
 - ▶ `ANY`
 - ▶ `EXISTS` / `NOT EXISTS`



The IN Operator

- ▶ Checks whether a column contains a value that is within the results of the subquery
- ▶ Example: Find all the courses that were offered both in Fall 2017 and Spring 2018

A screenshot of MySQL Workbench showing a query editor and a result grid. The query uses the IN operator to find course IDs from the section table that appear in both Fall 2017 and Spring 2018. The result grid shows one row with course_id 'CS-101'.

```
1 • SELECT DISTINCT course_id
2   FROM section
3   WHERE semester = 'Fall' AND year = 2017
4   AND course_id IN (
5     SELECT course_id
6     FROM section
7     WHERE semester = 'Spring' AND year = 2018
8   );
```

course_id
CS-101



The IN Operator

- ▶ You can compare more than one column with the results of the subquery
- ▶ Example: Find the total number of (distinct) students who have been taught by Katz

The screenshot shows a MySQL Workbench interface. The query editor contains the following SQL code:

```
1 • SELECT COUNT(DISTINCT id)
2   FROM takes
3   WHERE (course_id, sec_id, semester, year) IN (
4       SELECT course_id, sec_id, semester, year
5         FROM teaches AS t
6       JOIN instructor AS i
7         WHERE i.name = 'Katz'
8   );
```

The results grid below the editor shows a single row of data:

COUNT(DISTINCT id)
12



The NOT IN Operator

- ▶ Checks if the column doesn't contain any value within the results of the subquery
- ▶ Example: Find all the courses that were offered in Fall 2017 but not in Spring 2018

The screenshot shows a MySQL Workbench interface. The query editor window contains the following SQL code:

```
1 • SELECT DISTINCT course_id
2   FROM section
3   WHERE semester = 'Fall' AND year = 2017
4   AND course_id NOT IN (
5     SELECT course_id
6     FROM section
7     WHERE semester = 'Spring' AND year = 2018
8   );
```

The result grid below the editor displays the following data:

course_id
CS-347
PHY-101



The ALL Operator

- ▶ The ALL operator returns true if **all** of the subquery results meet the condition
- ▶ Used in conjunction with one of the comparison operators (=, <>, >, >=, <, or <=)
- ▶ Example: Find the department with the highest number of courses

A screenshot of MySQL Workbench showing a complex SQL query. The interface includes a toolbar at the top with various icons, a code editor with numbered lines 1 through 8, and a result grid below it.

```
1 •  SELECT dept_name, COUNT(*) AS courses_num
2   FROM course
3   GROUP BY dept_name
4   HAVING COUNT(*) >= ALL (
5       SELECT COUNT(*)
6       FROM course
7       GROUP BY dept_name
8   );
```

The result grid shows the output of the query:

dept_name	courses_num
Comp. Sci.	5



The ALL Operator

- ▶ Note that the following attempt to write the same query won't work

The screenshot shows a database query editor with the following code:

```
1 •  SELECT dept_name, COUNT(*) AS courses_num
2   FROM course
3   GROUP BY dept_name
4   HAVING COUNT(*) = (
5       SELECT MAX(COUNT(*))
6       FROM course
7       GROUP BY dept_name
8   );
```

The code attempts to use the ALL operator (indicated by the circled '•' symbol) in a HAVING clause. Below the code, the output pane shows an error message:

Output :::::
Action Output

#	Time	Action	Message
1	02:35:50	SELECT dept_name, COUNT(*) AS courses_num FROM course GROUP BY dept_name HAVING CO...	Error Code: 1111. Invalid use of group function

- ▶ You cannot layer aggregate functions on top of each other in the same SELECT



The ANY Operator

- ▶ The ANY operator returns true if **any** of the subquery results meets the condition
- ▶ Used in conjunction with one of the comparison operators (=, <>, >, >=, <, or <=)
- ▶ Example: Find all the instructors whose salary is higher than any of the instructors in the Biology department

The screenshot shows a database query interface with the following details:

Query Editor (Top):

```
1 • SELECT name, salary, dept_name
2   FROM instructor
3   WHERE salary > ANY (
4     SELECT salary
5     FROM instructor
6     WHERE dept_name = 'Biology'
7   );
```

Result Grid (Bottom):

	name	salary	dept_name
▶	Wu	90000.00	Finance
	Einstein	95000.00	Physics
	Gold	87000.00	Physics
	Katz	75000.00	Comp. Sci.
	Singh	80000.00	Finance
	Brandt	92000.00	Comp. Sci.
	Kim	80000.00	Elec. Eng.



The EXISTS Operator

- ▶ The EXISTS operator tests the existence of any result in the subquery
- ▶ It returns true if the subquery returns one or more rows
- ▶ For example, another way to write the query “Find all the courses that were offered both in Fall 2017 and Spring 2018”

A screenshot of MySQL Workbench showing a query editor window. The query is:

```
1 •  SELECT DISTINCT course_id
2   FROM section AS s1
3   WHERE semester = 'Fall' AND year = 2017
4   AND EXISTS (
5       SELECT *
6       FROM section AS s2
7       WHERE semester = 'Spring' AND year = 2018
8       AND s1.course_id = s2.course_id
9   );
```

The results pane shows a single row:

course_id
CS-101



The NOT EXISTS Operator

- ▶ You can also use NOT EXISTS to check if the subquery didn't return any rows
- ▶ Example: Find instructors who haven't taught any course

A screenshot of MySQL Workbench showing a query and its results. The query is:

```
1 • SELECT id, name
2   FROM instructor AS i
3 WHERE NOT EXISTS (
4     SELECT *
5       FROM teaches AS t
6      WHERE t.id = i.id
7 );
```

The results grid shows the following data:

	id	name
▶	33456	Gold
	58583	Califieri
	76543	Singh



The NOT EXISTS Operator

- ▶ Example: Find students who have taken all the courses in Music

The screenshot shows a database query editor window. At the top, there's a toolbar with various icons. Below it is a code editor pane containing the following SQL query:

```
1 •  SELECT id, name
2   FROM student AS s
3   WHERE NOT EXISTS (
4     SELECT *
5       FROM course AS c
6      WHERE dept_name = 'Music'
7     AND NOT EXISTS (
8       SELECT *
9         FROM takes AS t
10        WHERE s.id = t.id
11        AND c.course_id = t.course_id
12     )
13   );
```

At the bottom of the editor, there are buttons for "Result Grid", "Filter Rows", "Export", and "Wrap Cell Content". A preview grid shows one row of results:

	id	name
▶	55739	Sanchez



The NOT EXISTS Operator

- ▶ Another way to write the same query:

The screenshot shows a database query editor interface with a toolbar at the top and a results grid at the bottom. The query itself is as follows:

```
1 •  SELECT s.id, s.name
2   FROM student AS s
3   JOIN takes AS t
4   ON s.id = t.id
5   JOIN course AS c
6   ON t.course_id = c.course_id
7   WHERE c.dept_name = 'Music'
8   GROUP BY t.id
9   HAVING COUNT(DISTINCT t.course_id) = (
10      SELECT COUNT(*)
11      FROM course
12      WHERE dept_name = 'Music'
13    )
```

The results grid shows one row of data:

	id	name
▶	55739	Sanchez



Subqueries in the FROM Clause

- ▶ SQL allows you to write subqueries in the FROM clause
 - ▶ In MySQL, the subquery must be given a name using the AS keyword
- ▶ The subquery results can be joined with other tables or other subqueries
- ▶ Example: Find the highest number of courses offered in any given department

A screenshot of MySQL Workbench showing a query editor and a results grid. The query editor displays the following SQL code:

```
1 •  SELECT MAX(courses_in_dept.courses_num)
2   ⊖ FROM (
3     SELECT COUNT(*) courses_num
4     FROM course
5     GROUP BY dept_name
6   ) AS courses_in_dept;
```

The results grid shows a single row with the value 5, corresponding to the maximum number of courses offered in any department.



Updates with Subqueries

- ▶ You can also use subqueries in UPDATE and DELETE statements
- ▶ Example: recompute and update the tot_cred value for each student

```
UPDATE student AS s
SET tot_cred = (
    SELECT SUM(c.credits)
    FROM takes AS t
    JOIN course AS c
    WHERE t.course_id = c.course_id
        AND s.id = t.id
        AND t.grade IS NOT NULL
        AND t.grade <> 'F'
);
```

- ▶ Sets tot_cred to NULL for students who have not taken any course
- ▶ Instead of SUM(credits) can use COALESCE(SUM(credits), 0)
 - ▶ This would return SUM(credits) if it is not NULL, and 0 otherwise



Common Table Expressions (CTEs)

- ▶ CTE is a named temporary result set that only exists for the duration of the query
- ▶ CTEs are useful for breaking down complex queries into simpler parts
- ▶ Basic syntax of CTE:

```
WITH cte_name (column_name1, column_name2, ...)  
AS (  
    -- Your SQL query here  
)  
-- Your main SQL statement using the CTE here
```



Common Table Expressions (CTEs)

- ▶ Example: find all the instructors that earn more than the average salary

The screenshot shows a database query interface with the following details:

- Query Editor:** Displays the SQL code for finding instructors earning more than the average salary using a Common Table Expression (CTE).

```
1 • WITH instructor_avg_salary AS (
2     SELECT AVG(salary) AS average_salary
3     FROM instructor
4 )
5     SELECT name, salary
6     FROM instructor i, instructor_avg_salary a
7     WHERE i.salary > a.average_salary;
```
- Result Grid:** Shows the output of the query, listing the names and salaries of instructors who earn more than the average salary.

name	salary
Wu	90000.00
Einstein	95000.00
Gold	87000.00
Katz	75000.00
Singh	80000.00
Brandt	92000.00
Kim	80000.00



Common Table Expressions (CTEs)

- ▶ Show the distribution of number of courses taken by students
 - ▶ i.e., how many students took only one course, how many students took two courses, etc.

The screenshot shows a database query editor interface. At the top, there's a toolbar with various icons. Below it, the SQL code is displayed:

```
1 • 1 WITH courses_per_student AS (
2     -- Count the number of courses each student took
3     SELECT id, COUNT(DISTINCT course_id) AS num_courses
4     FROM takes
5     GROUP BY id
6 )
7
8     -- Get the distribution of courses taken by students
9     SELECT num_courses, COUNT(id) AS num_students
10    FROM courses_per_student
11    GROUP BY num_courses
12    ORDER BY num_courses;
```

At the bottom, a result grid shows the output:

	num_courses	num_students
▶	1	5
	2	6
	4	1



Recursive Queries

- ▶ A recursive query is a query that refers back to itself
- ▶ Used to deal with hierarchical or tree-structured data
- ▶ A recursive query is written using a CTE that consists of two subqueries:
 - ▶ The base (non-recursive) query provides the starting point of the recursion
 - ▶ The recursive query which has a reference to the CTE itself



Recursive Queries

- ▶ Example: Find for each course all its prerequisites (both direct and indirect)

```
1 • WITH RECURSIVE rec_prereq(course_id, prereq_id) AS (
  2   -- Anchor member: The base case, where we select the initial prerequisites
  3   SELECT course_id, prereq_id
  4   FROM prereq
  5   UNION
  6   -- Recursive member: Select the prerequisites of the previous rows
  7   SELECT p.course_id, p.prereq_id
  8   FROM prereq AS p
  9   JOIN rec_prereq AS rec_p
 10  ON p.course_id = rec_p.prereq_id -- Join with the CTE itself!
 11 )
 12 -- Now, select the results from the CTE
 13 SELECT * FROM rec_prereq;
```

< Result Grid | Filter Rows: | Export: | Wrap Cell Content:

course_id	prereq_id
BIO-301	BIO-101
BIO-399	BIO-101
CS-190	CS-101
CS-315	CS-101
CS-319	CS-101
CS-347	CS-101
EE-181	PHY-101

Summary: SQL Execution Order



Northeastern
University

1. FROM
2. ON
3. JOIN
4. WHERE
5. GROUP BY
6. WITH CUBE or WITH
ROLLUP
7. HAVING
8. SELECT
9. DISTINCT
10. ORDER BY
11. TOP





Query Execution Plans

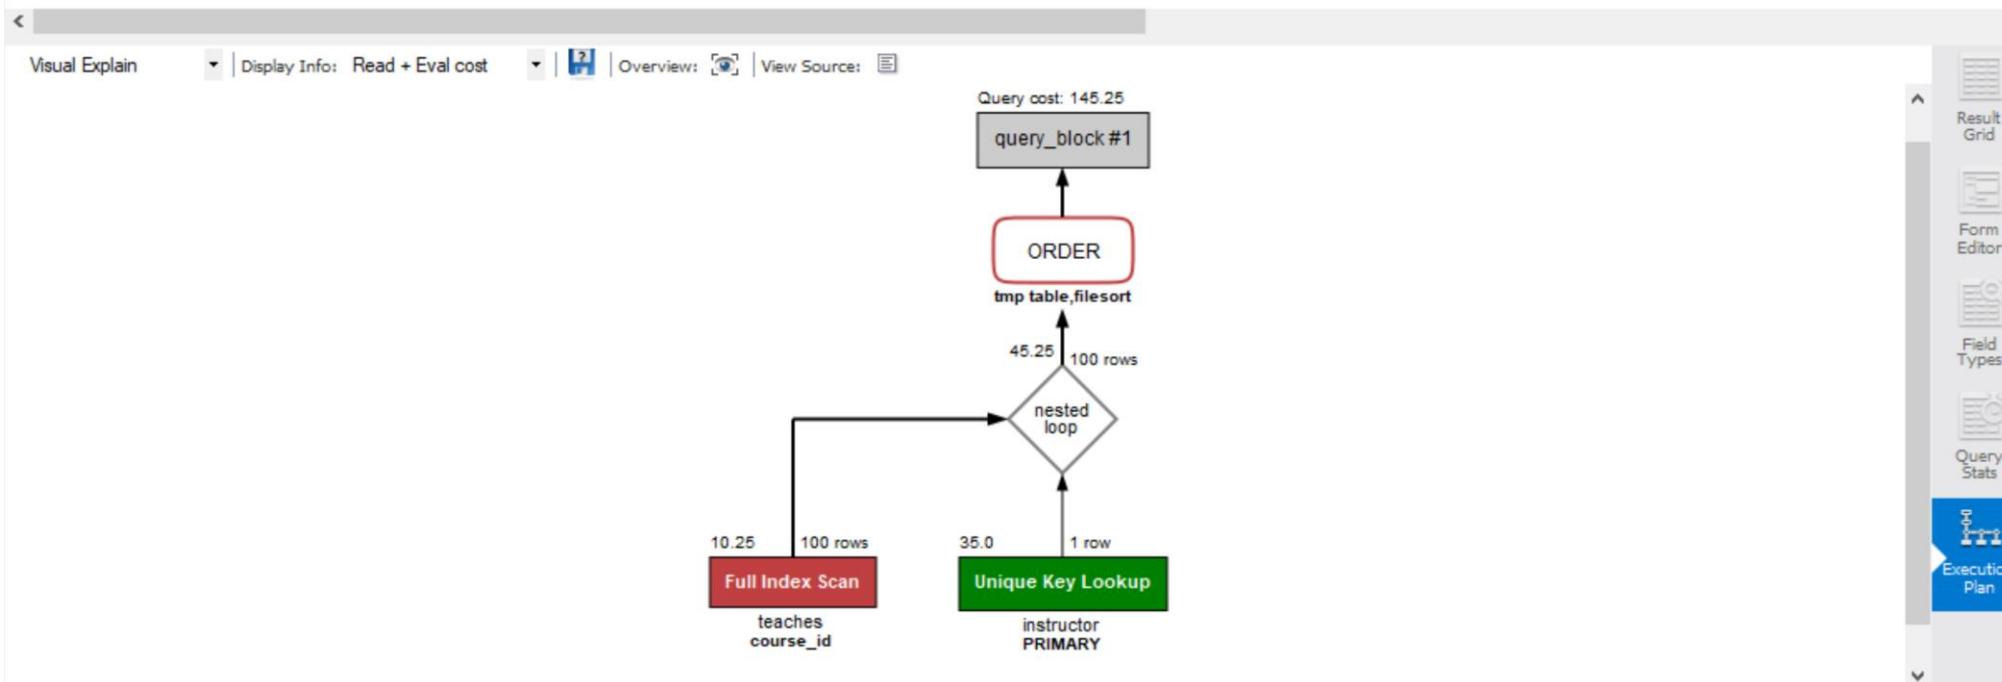
- ▶ A **query execution plan** is an ordered set of steps used to access data in the database
 - ▶ Defines which algorithm to use for each operation, and how to coordinate the operations
- ▶ When a query is submitted to the database, the **query optimizer** evaluates the different plans for executing the query and returns what it considers the best option
 - ▶ Cost difference between execution plans for a query can be enormous
- ▶ Estimation of plan cost is based on:
 - ▶ Statistical information about tables
 - ▶ e.g., number of rows, number of distinct values in a column
 - ▶ Cost of different algorithms (e.g., join vs. subquery)
 - ▶ Also computed using statistics
- ▶ Sometimes we need to manually examine and tune the optimizer plans



Query Execution Plans

- To view the execution plan, click on Execution Plan next to the query results

```
1 • SELECT *
2   FROM instructor
3   JOIN teaches
4     ON instructor.id = teaches.id
5   ORDER BY name;
```





Views

- ▶ A view is a virtual table based on the results set of an SQL statement
- ▶ Reasons to use views:
 - ▶ Data security – allows you to hide certain data from view of certain users
 - ▶ e.g., some users only need to know the instructor IDs and names, but not their salaries
 - ▶ Hiding complexity of the underlying tables
- ▶ A view is created with the CREATE VIEW statement:

```
CREATE VIEW view_name AS
SELECT column1, column2, ...
FROM table_name
WHERE condition;
```

- ▶ A view always shows up-to-date data
 - ▶ The database engine recreates the view whenever it is used in a query



Views

- ▶ Example: A view that shows all Physics instructors without their salary

```
CREATE VIEW physics_faculty AS
SELECT id, name
FROM instructor
WHERE dept_name = 'Physics';
```

- ▶ We can query the view above as follows:

The screenshot shows a database management system interface. At the top, there is a toolbar with various icons. Below the toolbar, the SQL query is displayed:

```
1 • SELECT *
2   FROM physics_faculty;
```

Below the query, the results are shown in a "Result Grid". The grid has two columns: "id" and "name". The data is as follows:

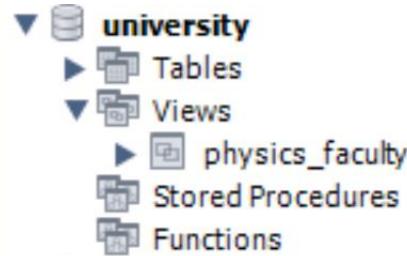
	id	name
▶	22222	Einstein
	33456	Gold

- ▶ View names may appear in a query anywhere a table name may appear



Views

- ▶ You can see all the available views in the database under the View folder:



- ▶ To update a view use the CREATE OR REPLACE VIEW command:

```
CREATE OR REPLACE VIEW view_name AS
SELECT column1, column2, ...
FROM table_name
WHERE condition;
```

- ▶ To delete a view use the DROP VIEW command:

```
DROP VIEW view_name;
```



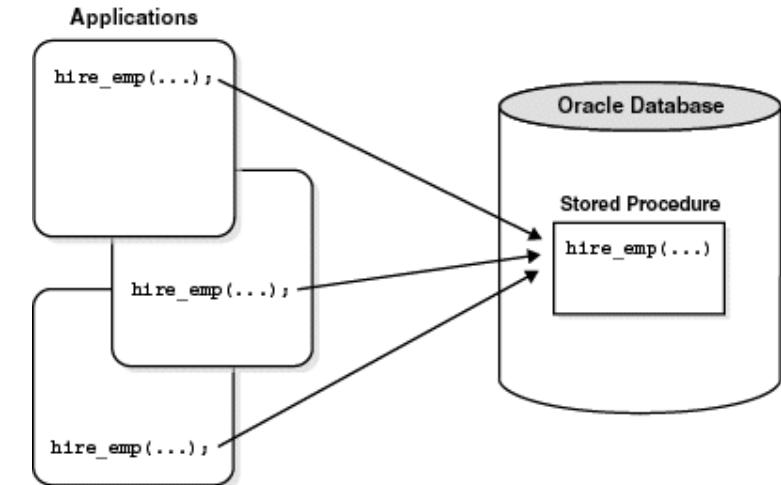
Updatable Views

- ▶ Some views are updatable
 - ▶ i.e., you can use them in statements like INSERT or UPDATE to update the underlying tables
- ▶ For a view to be updatable, there must be a one-to-one relationship between the rows in the view and the rows in the underlying table
- ▶ More specifically, the view isn't updatable if it contains any of the following:
 - ▶ Aggregate functions like MIN() or AVG()
 - ▶ DISTINCT
 - ▶ GROUP BY or HAVING
 - ▶ Certain joins are not allowed
 - ▶ Subquery in SELECT
 - ▶ Subquery in the WHERE clause that refers to a table in the FROM clause
 - ▶ UNION or UNION ALL



Stored Procedures

- ▶ A **stored procedure (SP)** is a group of SQL statements that is stored in the database
 - ▶ You can also pass parameters to a stored procedure
- ▶ The benefits of using stored procedures:
 - ▶ Reusable code that can be used by different applications
 - ▶ Allows a single point of change in case the business rules change
 - ▶ Faster execution
 - ▶ Reduce the network traffic
 - ▶ More secure than ad-hoc queries





Creating a Stored Procedure

- ▶ The syntax for creating a new stored procedure in MySQL:

```
DELIMITER $$  
  
CREATE PROCEDURE procedure_name (  
    IN | OUT | INOUT parameter1 datatype,  
    IN | OUT | INOUT parameter2 datatype,  
    ...  
)  
BEGIN  
    -- SQL statements  
END $$  
DELIMITER ;
```

Redefine the delimiter temporarily so that ; can be used to separate statements inside the procedure

Use ; in the body of the stored procedure

Use the delimiter \$\$ to end the stored procedure

Change the default delimiter back to ;

- ▶ Parameter modes:

- ▶ **in** (default) – the parameter's value is passed into the SP and cannot be changed inside the SP
- ▶ **out** – the parameter's value is passed back to the calling program, must be a variable
- ▶ **inout** – the parameter's value is passed to the SP and a new value can be assigned to it



Creating a Stored Procedure

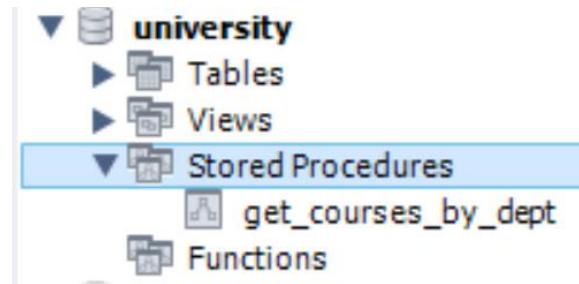
- ▶ Example: a stored procedure that returns all courses in a given department

```
1  DELIMITER $$  
2  
3 • ⊕ CREATE PROCEDURE get_courses_by_dept(  
4     IN dept_name VARCHAR(20)  
5 )  
6 ⊕ BEGIN  
7     SELECT *  
8     FROM course AS c  
9     WHERE c.dept_name = dept_name;  
10    END $$  
11  
12    DELIMITER ;
```



Creating a Stored Procedure

- ▶ You can view your procedure under the Stored Procedures folder of the schema





Calling a Stored Procedure

- To execute the store procedure, you use the CALL keyword:

```
CALL procedure_name(parameter1, parameter2, ...);
```

- For example, let's call the get_courses_by_dept with Comp. Sci. as the argument:

```
1 CALL get_courses_by_dept('Comp. Sci.');
```

course_id	title	dept_name	credits
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3



Calling a Stored Procedure

- You can also execute the SP by clicking on the Execute button next to its name

The screenshot shows the MySQL Workbench interface. On the left, the 'university' database is selected, and the 'Stored Procedures' node is expanded, showing the 'get_courses_by_dept' procedure. In the center, a dialog box titled 'Call stored procedure university.get_courses_by_dept' is open, prompting for parameter values. The parameter 'dept_name' is set to 'Comp. Sci.' with the type 'VARCHAR(20)'. At the bottom of the dialog are 'Execute' and 'Cancel' buttons. Below the dialog, the SQL editor window displays the following code:

```
1 • call university.get_courses_by_dept('Comp. Sci.');?>
2
```

At the bottom, the 'Result Grid' shows the results of the executed query:

course_id	title	dept_name	credits
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3



OUT Parameter

- ▶ A stored procedure that returns the number of students in a given department

```
1  DELIMITER $$  
2  
3 •◦ CREATE PROCEDURE get_enrollment_by_dept(  
4      IN dept_name VARCHAR(20),  
5      OUT students_num INT  
6  )  
7  BEGIN  
8      SELECT COUNT(*) INTO students_num  
9      FROM student AS s  
10     WHERE s.dept_name = dept_name;  
11  END $$  
12  
13  DELIMITER ;
```



OUT Parameter

- ▶ To call a SP with an OUT parameter you need to pass a **session variable** to it
 - ▶ This variable will receive the return value from the procedure
- ▶ Session variables are variables that start with @
 - ▶ Don't require any declaration
 - ▶ Can be used inside any SQL query or statement
 - ▶ Exist until the end of the current session
 - ▶ Assignments to the variable are performed using a SET statement
 - ▶ To display the value of a variable use the SELECT statement

The screenshot shows a MySQL Workbench interface. At the top, there's a toolbar with various icons. Below the toolbar, two lines of SQL code are displayed:

```
1 • CALL get_enrollment_by_dept('Comp. Sci.', @num);
2 • SELECT @num AS students_num;
```

Below the code, there's a "Result Grid" section. It has a header row with a single column labeled "students_num". Underneath is a data row containing the value "4".



InOut Parameter Example

- ▶ The following example demonstrates how to use an **inout** parameter in a SP:

```
1 delimiter $$  
2 •◦ create procedure update_counter (  
3     inout counter int,  
4     in increment int  
5 )  
6 begin  
7     set counter = counter + increment;  
8 end $$  
9  
10 delimiter ;
```



InOut Parameter Example

- ▶ The following statements illustrate how to call the update_counter SP:

The screenshot shows a MySQL Workbench interface. At the top, there is a toolbar with various icons. Below it, a text area contains the following SQL code:

```
1 • set @counter = 1;
2 • call update_counter(@counter, 1); -- 2
3 • call update_counter(@counter, 1); -- 3
4 • call update_counter(@counter, 5); -- 8
5 • select @counter; -- 8
6
7
```

At the bottom, a result grid displays the value of the variable @counter:

@counter
8



Programming Language Constructs

- ▶ Stored procedures support programming language constructs like ifs and loops:

```
IF boolean expression THEN
    statement_list
[ELSEIF boolean expression THEN
    statement_list]
...
[ELSE statement_list]
END IF;
```

```
WHILE boolean expression DO
    statement_list
END WHILE;
```

```
REPEAT
    statement_list
UNTIL boolean expression
END REPEAT;
```



Programming Language Constructs

- The following procedure registers a student to a course section only if there is enough room left in the classroom allocated for this section

```
1  DELIMITER $$  
2  
3 •◦ CREATE PROCEDURE register_student(  
4     IN student_id VARCHAR(5),  
5     IN course_id VARCHAR(8),  
6     IN sec_id VARCHAR(8),  
7     IN semester VARCHAR(6),  
8     IN year NUMERIC(4, 0),  
9     OUT success INT,  
10    OUT error_msg VARCHAR(100)  
11 )  
12 ◦ BEGIN  
13     DECLARE curr_enrollment INT;  
14     DECLARE classroom_capacity INT;  
15  
16     SELECT COUNT(*) INTO curr_enrollment  
17     FROM takes AS t  
18     WHERE t.course_id = course_id AND t.sec_id = sec_id  
19         AND t.semester = semester AND t.year = year;
```

Programming Language Constructs



```
21  SELECT capacity INTO classroom_capacity
22  FROM classroom AS c
23  JOIN section AS s
24  ON c.building = s.building AND c.room_number = s.room_number
25  WHERE s.course_id = course_id AND s.sec_id = sec_id
26      AND s.semester = semester AND s.year = year;
27
28  IF curr_enrollment < classroom_capacity THEN
29      INSERT INTO takes
30      VALUES (student_id, course_id, sec_id, semester, year, NULL);
31      SET success = 1;
32  ELSE
33      SET success = 0;
34  SET error_msg = CONCAT('Enrollment limit reached for course ', course_id,
35      ' section ', sec_id);
36  END IF;
37
38
39  DELIMITER ;
```

Programming Language Constructs



Northeastern
University

- ▶ Calling the procedure:

The screenshot shows a MySQL Workbench interface. At the top, there is a toolbar with various icons. Below the toolbar, a query editor window contains the following SQL code:

```
1 CALL register_student(98988, 'CS-101', 1, 'Fall', 2017, @success, @error_msg);
2 • SELECT @success, @error_msg;
```

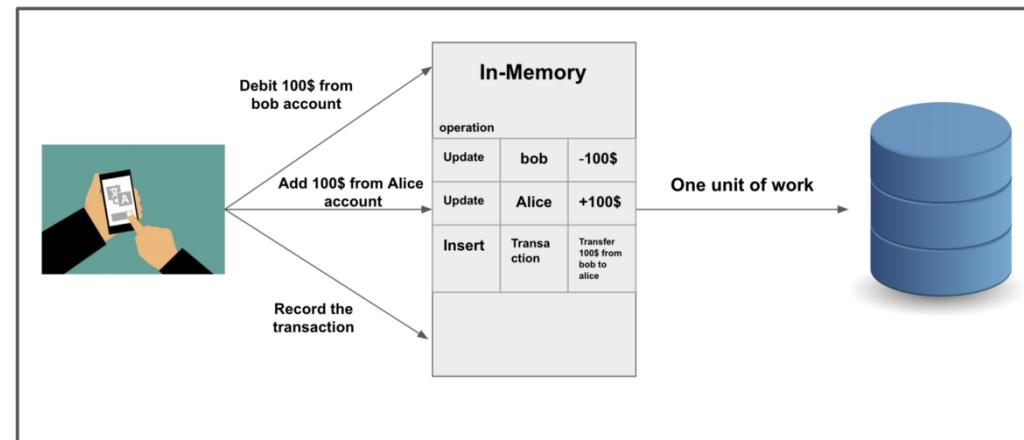
Below the query editor is a results grid titled "Result Grid". The grid has two columns: "@success" and "@error_msg". A single row of data is shown, with the "@success" value being "1" and the "@error_msg" value being "NULL".

	@success	@error_msg
▶	1	NULL



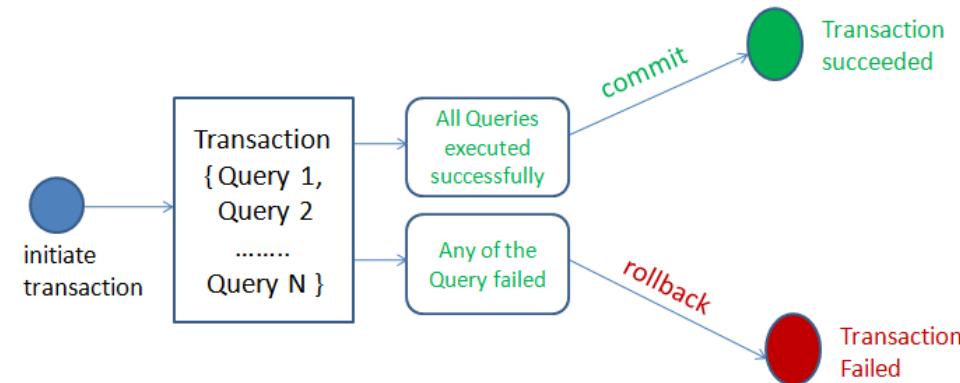
Transactions

- ▶ A transaction is a sequence of SQL statements that represents a **single unit of work**
- ▶ For example, consider a Bank amount transfer, which involves two operations:
 - ▶ Withdrawal of money from account *A*
 - ▶ Deposit money to Account *B*
- ▶ If the system crashes after subtracting the amount from *A* but before adding it to *B*, the bank balances will be inconsistent



Transactions

- ▶ To define a new transaction, use the START TRANSACTION statement
- ▶ All the subsequent SQL statements will belong to the transaction
- ▶ The transaction must end with one of the following statements:
 - ▶ COMMIT: The updates performed by the transaction become permanent in the database
 - ▶ ROLLBACK: All updates performed by the transaction are undone





Transactions

- ▶ In MySQL, transactions need to be written inside a stored procedure
 - ▶ So you can automatically rollback the transaction whenever an error occurs
 - ▶ The following statement needs to be written in the beginning of the procedure:

```
DECLARE EXIT HANDLER FOR SQLEXCEPTION ROLLBACK;
```
- ▶ The following procedure updates a student's grade in a course
- ▶ This procedure needs to update both the takes table and the student's tot_cred
 - ▶ These two updates should be performed in a single transaction



Transaction Example

```
1  DELIMITER $$  
2  
3 • CREATE PROCEDURE update_grade(  
4      student_id VARCHAR(5),  
5      course_id VARCHAR(8),  
6      sec_id VARCHAR(8),  
7      semester VARCHAR(6),  
8      year DECIMAL(4),  
9      grade VARCHAR(2)  
10 )  
11 BEGIN  
12     -- If an error occurs, the transaction will be rolled back automatically  
13     DECLARE EXIT HANDLER FOR SQLEXCEPTION ROLLBACK;  
14  
15     START TRANSACTION;  
16     -- Update the student's grade in the course  
17     UPDATE takes AS t  
18     SET t.grade = grade  
19     WHERE t.id = student_id AND t.course_id = course_id AND t.sec_id = sec_id  
20         AND t.semester = semester AND t.year = year;
```



Transaction Example

```
22    -- Update the student's total credits
23    UPDATE student AS s
24    SET s.tot_cred = s.tot_cred + (
25        SELECT credits
26        FROM course AS c
27        WHERE c.course_id = course_id
28    )
29    WHERE s.id = student_id;
30
31    COMMIT;
32 END $$

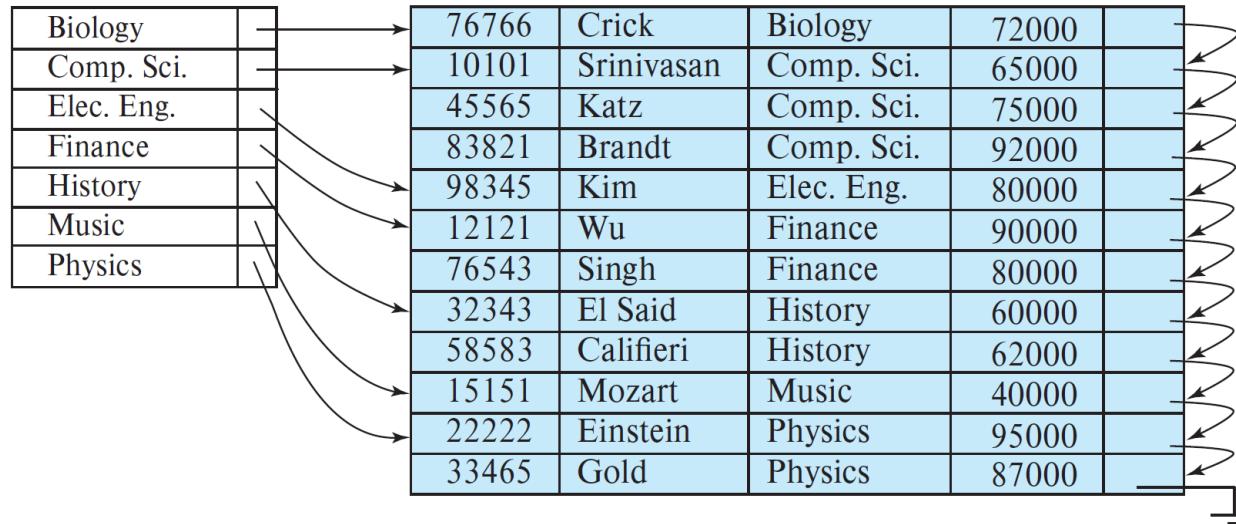
33

34 DELIMITER ;
```



Indexes

- ▶ An index is a data structure that can speed up queries/searches
 - ▶ e.g., searching for instructors in a specific department



- ▶ Drawbacks of indexes
 - ▶ Causes updates to the table to become slower (since the index also needs to be updated)
 - ▶ Storage space

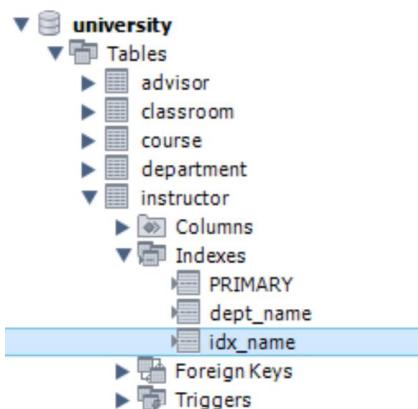


Creating Indexes

- ▶ We create an index using the CREATE INDEX command

```
CREATE [UNIQUE] INDEX index_name  
ON table_name (column1, column2, ...);
```

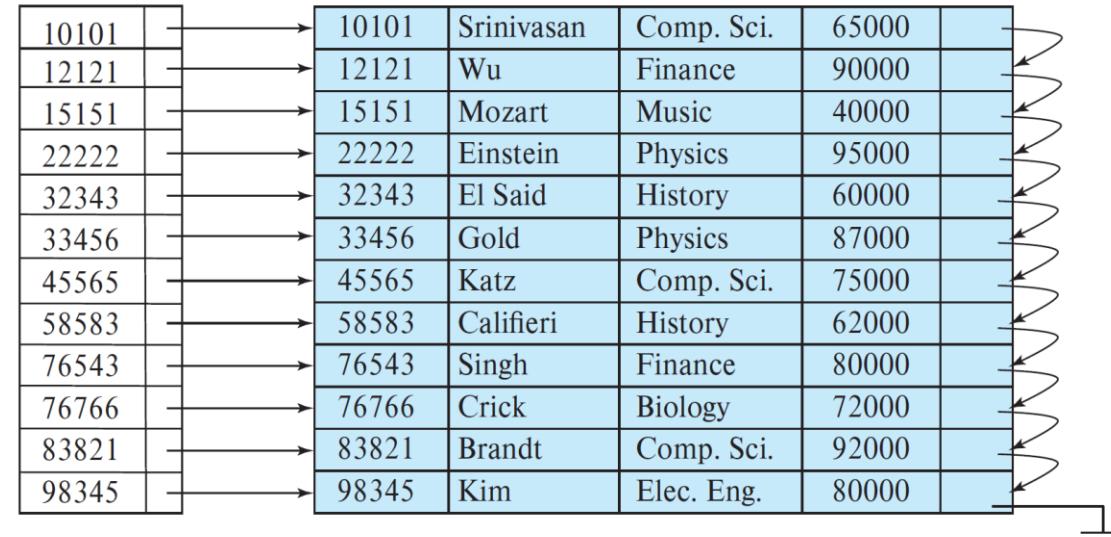
- ▶ A UNIQUE index creates a constraint such that all values in the index must be distinct
- ▶ For example, let's define an index on the name column in the instructor table:
- ▶ You can see the new index under the Indexes folder of the table:





Clustered Indexes

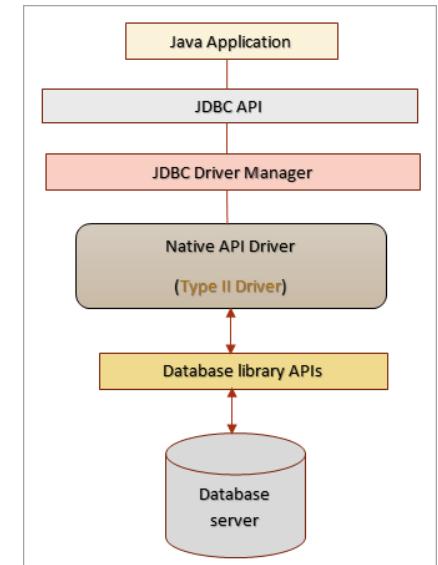
- ▶ A **clustered index** defines the physical order in which table records are stored
- ▶ There can be only one clustered index per table
 - ▶ By default a clustered index is created on a primary key column
- ▶ Accessing a row through the clustered index is fast
 - ▶ Since the index search leads directly to the page that contains the row data





Accessing SQL from a Programming Languages

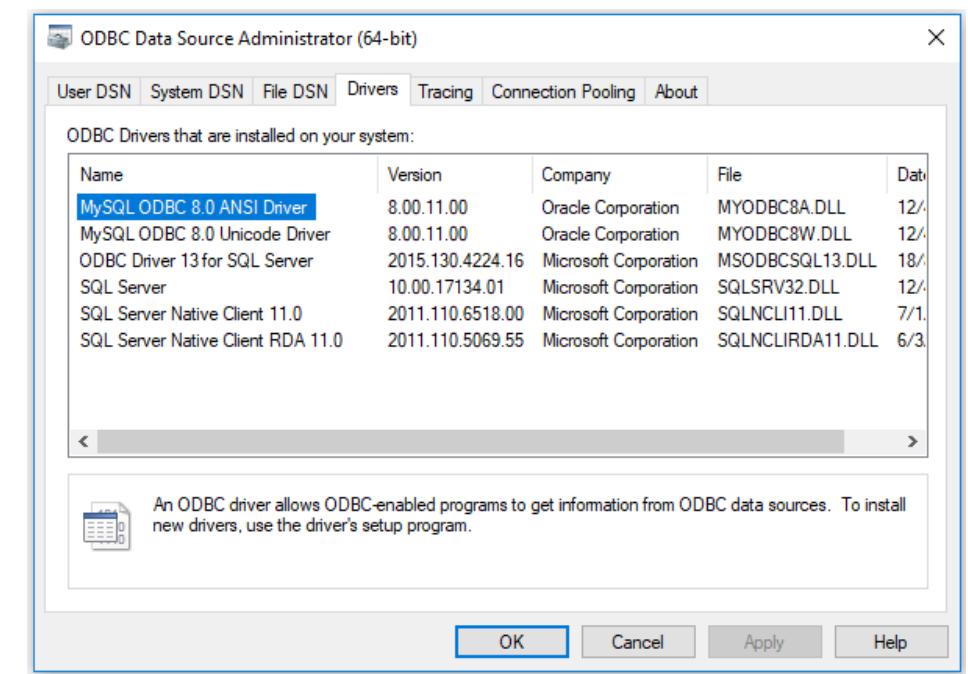
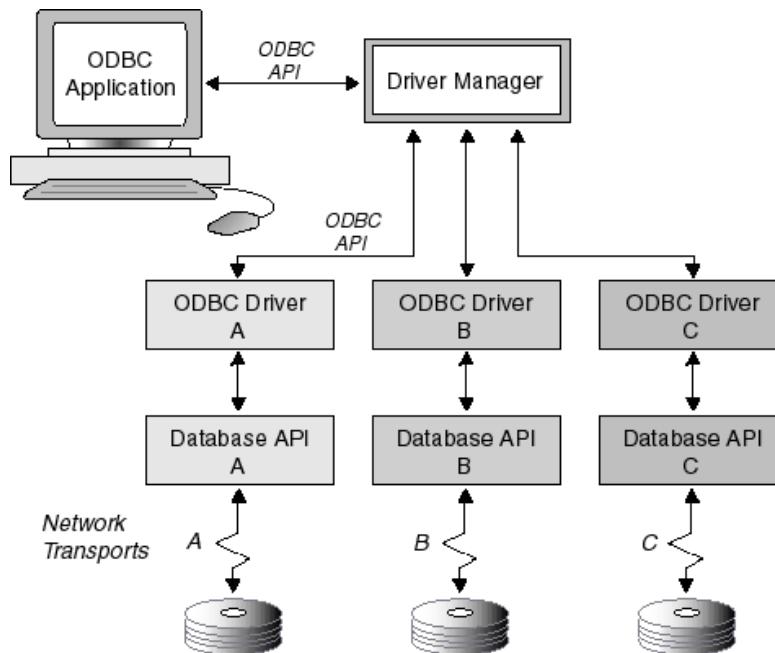
- ▶ To access SQL from a programming language you need a **database driver**
- ▶ A driver converts application commands into a format that the DBMS can understand
- ▶ Typical commands include
 - ▶ Connection commands (e.g., open or close a connection to the database)
 - ▶ SQL commands
 - ▶ Handling cursors, which allow you to traverse through result sets
 - ▶ Transaction management commands
- ▶ Types of drivers
 - ▶ Native drivers – provided by specific database vendors
 - ▶ ODBC drivers – provide a standard interface to communicate with databases
 - ▶ JDBC drivers – specific to Java-based applications



ODBC



- ▶ Open Database Connectivity (ODBC) is a standard API for accessing DBMS
- ▶ Independent of any specific DBMS or operating system
- ▶ ODBC drivers exist for most of the commercial databases



MySQL Connector/Python

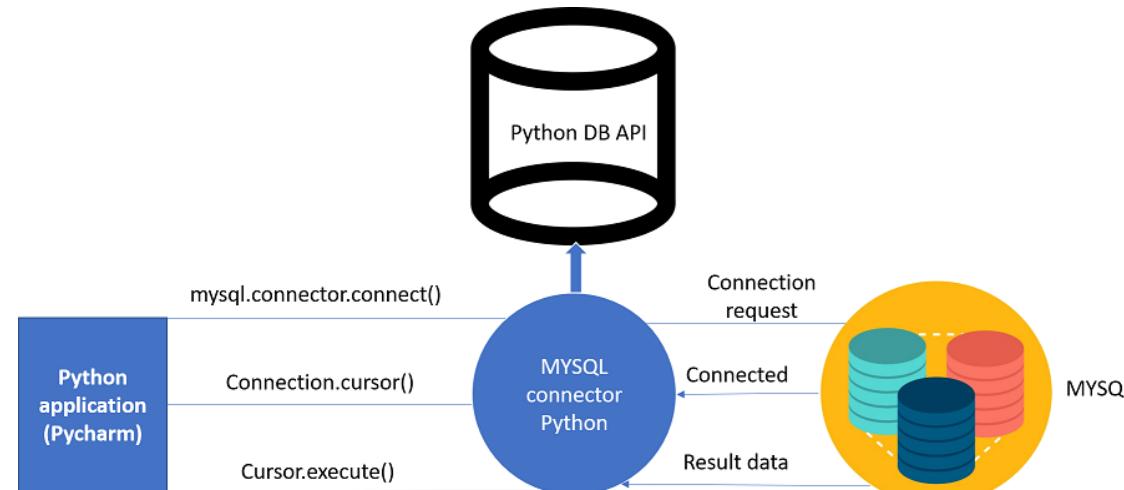


- ▶ **Python Database API (DB-API)** defines a standard interface for accessing relational databases from Python programs
- ▶ Different packages in Python implement this interface for different databases
 - ▶ e.g., mysql.connector for MySQL, sqlite3 for SQLite, pyodbc for ODBC
- ▶ mysql.connector is the recommended driver for interacting with MySQL
 - ▶ Developed by the MySQL group at Oracle
- ▶ Install it via pip
 - ```
pip install mysql-connector-python
```
  - ▶ If you've installed Python with the Anaconda distribution, you should already have it



# MySQL Connector/Python

- ▶ The general workflow of a Python program that interacts with a MySQL database:
  - ▶ Connect to the MySQL server
  - ▶ Execute a SQL query
  - ▶ Fetch the results
  - ▶ Inform the database if any table changes were made (by committing the changes)
  - ▶ Close the connection to the server





# Establishing a Connection with MySQL Server

- ▶ To connect to a MySQL server, call the **connect()** function in `mysql.connector` module
  - ▶ This function gets 4 parameters: host, user, password and database name
  - ▶ It returns a MySQLConnection object

```
from mysql.connector import connect, Error
from getpass import getpass
```

```
try:
 conn = connect(
 host='localhost',
 user=input('Enter username: '),
 password=getpass('Enter password: '),
 database='university'
)
except Error as e:
 print(e)
```

- ▶ You should always close the connection in the end by calling **conn.close()**
- ▶ You should never hard-code your login credentials directly in a Python script



# The Cursor Object

- ▶ In order to execute SQL queries in Python, you need a **cursor** object
  - ▶ A cursor object allows you to traverse over database records
- ▶ To create a cursor, use the **cursor()** method of your connection object
- ▶ Then you execute a SQL query by calling **cursor.execute(query)**
- ▶ If the query returns rows, you can retrieve them using one of cursor's fetch methods:
  - ▶ **fetchall()** - retrieves all the rows from the result as a list of tuples
  - ▶ **fetchone()** - retrieves the next row of the result as a tuple
    - ▶ Returns None if no more rows are available
  - ▶ **fetchmany(*n*)** - retrieves the next *n* rows from the result as a list of tuples (*n* defaults to 1)
    - ▶ Returns an empty list if no more rows are available



# Reading Records from a Table

- The following example selects all the records from the *instructor* table:

```
inst_query = "SELECT * FROM instructor"
with conn.cursor() as cursor:
 cursor.execute(inst_query)
 result = cursor.fetchall()
 for row in result:
 print(row)

('10101', 'Srinivasan', 'Comp. Sci.', Decimal('65000.00'))
('12121', 'Wu', 'Finance', Decimal('90000.00'))
('15151', 'Mozart', 'Music', Decimal('40000.00'))
('22222', 'Einstein', 'Physics', Decimal('95000.00'))
('32343', 'El Said', 'History', Decimal('60000.00'))
('33456', 'Gold', 'Physics', Decimal('87000.00'))
('45565', 'Katz', 'Comp. Sci.', Decimal('75000.00'))
('58583', 'Califieri', 'History', Decimal('62000.00'))
('76543', 'Singh', 'Finance', Decimal('80000.00'))
('76766', 'Crick', 'Biology', Decimal('72000.00'))
('83821', 'Brandt', 'Comp. Sci.', Decimal('92000.00'))
('98345', 'Kim', 'Elec. Eng.', Decimal('80000.00'))
```

- The result variable holds the records returned from `fetchall()`
- It's a list of tuples representing individual records from the table



# Reading Records from a Table

- ▶ To get specific attributes in each row, specify their indexes in the returned tuples
- ▶ For example, to print only the instructor names:

```
inst_query = "SELECT * FROM instructor"
with conn.cursor() as cursor:
 cursor.execute(inst_query)
 result = cursor.fetchall()
 for row in result:
 print(row[1])
```

Srinivasan  
Wu  
Mozart  
Einstein  
El Said  
Gold  
Katz  
Califieri  
Singh  
Crick  
Brandt  
Kim



# Cursor as an Iterator

- To process the rows in the result one at a time, you can use the cursor as an iterator:

```
inst_query = "SELECT * FROM instructor"
with conn.cursor() as cursor:
 cursor.execute(inst_query)
 for row in cursor:
 print(row)

('10101', 'Srinivasan', 'Comp. Sci.', Decimal('65000.00'))
('12121', 'Wu', 'Finance', Decimal('90000.00'))
('15151', 'Mozart', 'Music', Decimal('40000.00'))
('22222', 'Einstein', 'Physics', Decimal('95000.00'))
('32343', 'El Said', 'History', Decimal('60000.00'))
('33456', 'Gold', 'Physics', Decimal('87000.00'))
('45565', 'Katz', 'Comp. Sci.', Decimal('75000.00'))
('58583', 'Califieri', 'History', Decimal('62000.00'))
('76543', 'Singh', 'Finance', Decimal('80000.00'))
('76766', 'Crick', 'Biology', Decimal('72000.00'))
('83821', 'Brandt', 'Comp. Sci.', Decimal('92000.00'))
('98345', 'Kim', 'Elec. Eng.', Decimal('80000.00'))
```



# More Complex Queries

- ▶ You can make your **select** queries as complex as you want using the same methods
- ▶ Example: Find the courses that were taken by the highest number of students

```
most_popular_courses = """
 SELECT C.course_id, C.title
 FROM course as C, takes as T
 WHERE C.course_id = T.course_id
 GROUP BY T.course_id
 HAVING count(*) >= ALL (
 SELECT count(*)
 FROM takes
 GROUP BY course_id
)
"""
with conn.cursor() as cursor:
 cursor.execute(most_popular_courses)
 result = cursor.fetchall()
 print(result)

[('CS-101', 'Intro. to Computer Science')]
```



# Cursor Properties

- ▶ The cursor has a few useful properties that provide information about the result set
  - ▶ **column\_names** – returns the list containing the column names of the result set
  - ▶ **rowcount** – the number of rows in the result set

```
inst_query = "SELECT * FROM instructor"
with conn.cursor() as cursor:
 cursor.execute(inst_query)
 cursor.fetchall()
 print('Column names:', cursor.column_names)
 print('Row count:', cursor.rowcount)
```

Column names: ('ID', 'name', 'dept\_name', 'salary')

Row count: 12



# SQL Injection Attack

- ▶ A common attack that enables execution of malicious SQL statements in the DB
  - ▶ By insertion (“injection”) of a SQL query via the input data from the client
- ▶ Suppose you write a script that checks if a given username exists in the users table
- ▶ You construct the following SQL query:

```
username = # read from the input
query = "SELECT * FROM users WHERE username = '" + username + "'"
```

- ▶ If the user, instead of entering their username, enters:

Username:

- ▶ Then the query becomes:

```
query = "SELECT * FROM users WHERE username = '' or 1 = 1 --'"
```

- ▶ The **where** clause is now always true and the entire users table is returned



# SQL Injection Attack

- ▶ Most databases allow execution of multiple SQL statements separated by semicolon
- ▶ This allows the hacker to inject whole SQL statements into the query
- ▶ For example, the hacker could enter the following string:

Username: '`;` drop table users; --

Login

- ▶ The resulting query would be:

```
query = "SELECT * FROM users WHERE username = ''; drop table users; --"
```

- ▶ This query would result in deleting the entire users table!
- ▶ **Solution:** use parameterized queries whenever user input is involved in the query



# SQL Injection Attack

## ▶ Example:

```
inst_name = input('Enter instructor name: ')
inst_query = f"""
 SELECT * FROM instructor
 WHERE name = '{inst_name}'
"""

print(inst_query)
with conn.cursor() as cursor:
 cursor.execute(inst_query)
 for row in cursor:
 print(row)
```

Enter instructor name: X' or 'Y' = 'Y

```
SELECT * FROM instructor
WHERE name = 'X' or 'Y' = 'Y'

('10101', 'Srinivasan', 'Comp. Sci.', Decimal('65000.00'))
('12121', 'Wu', 'Finance', Decimal('90000.00'))
('15151', 'Mozart', 'Music', Decimal('40000.00'))
('22222', 'Einstein', 'Physics', Decimal('95000.00'))
('32343', 'El Said', 'History', Decimal('60000.00'))
('33456', 'Gold', 'Physics', Decimal('87000.00'))
('45565', 'Katz', 'Comp. Sci.', Decimal('75000.00'))
```



# Parameterized Query

- ▶ A **parameterized query** is a query that uses placeholders (%) for attribute values
- ▶ Strings passed to the placeholders are correctly escaped by the library at runtime
  - ▶ e.g., each quotation mark is doubled
- ▶ You can pass a parameterized query to **cursor.execute()** as follows:

```
inst_name = input('Enter instructor name: ')
inst_query = """
 SELECT * FROM instructor
 WHERE name = %s
"""

val_tuple = (inst_name,)

with conn.cursor() as cursor:
 cursor.execute(inst_query, val_tuple)
 for row in cursor:
 print(row)
```

```
Enter instructor name: Katz
('45565', 'Katz', 'Comp. Sci.', Decimal('75000.00'))
```



# Parameterized Query

- ▶ If a user tries to sneak in some problematic characters, the resulting statement will cause no harm since each quotation mark will be doubled:

```
select * from instructor where name = 'X" or "Y" = "Y'
```

- ▶ This query will return an empty set

```
inst_name = input('Enter instructor name: ')
inst_query = """
 SELECT * FROM instructor
 WHERE name = %s
"""

val_tuple = (inst_name,)

with conn.cursor() as cursor:
 cursor.execute(inst_query, val_tuple)
 for row in cursor:
 print(row)
```

Enter instructor name: X' or 'Y' = 'Y



# Inserting New Records

- ▶ To insert data, pass the **insert into** command to the cursor's **execute()** method
- ▶ For example, to add a new instructor:

```
add_inst_query = """
 INSERT INTO instructor
 VALUES ('99987', 'Chris', 'Comp. Sci.', 98000)
"""

with conn.cursor() as cursor:
 cursor.execute(add_inst_query)
 conn.commit()
```

- ▶ You must call **conn.commit()** at the end, otherwise your changes will be lost!
  - ▶ Unless you turn on automatic commits by setting **conn.autocommit = True**



# Inserting New Records

- ▶ Verifying that the new record was added to the table:

The screenshot shows the MySQL Workbench interface with a query editor and a result grid. The query editor contains the SQL command: `1 • SELECT * FROM university.instructor;`. The result grid displays 14 rows of data from the `instructor` table, including columns `ID`, `name`, `dept_name`, and `salary`. The last row is a blank entry with all columns set to `NULL`.

|   | ID    | name       | dept_name  | salary   |
|---|-------|------------|------------|----------|
| ▶ | 10101 | Srinivasan | Comp. Sci. | 65000.00 |
|   | 12121 | Wu         | Finance    | 90000.00 |
|   | 15151 | Mozart     | Music      | 40000.00 |
|   | 22222 | Einstein   | Physics    | 95000.00 |
|   | 32343 | El Said    | History    | 60000.00 |
|   | 33456 | Gold       | Physics    | 87000.00 |
|   | 45565 | Katz       | Comp. Sci. | 75000.00 |
|   | 58583 | Califieri  | History    | 62000.00 |
|   | 76543 | Singh      | Finance    | 80000.00 |
|   | 76766 | Crick      | Biology    | 72000.00 |
|   | 83821 | Brandt     | Comp. Sci. | 92000.00 |
|   | 98345 | Kim        | Elec. Eng. | 80000.00 |
| * | 99987 | Chris      | Comp. Sci. | 98000.00 |
| * | NULL  | NULL       | NULL       | NULL     |



# Inserting New Records

- If the values to be inserted come from an external source (e.g., the user), again you should use parameters in the SQL statement:

```
inst_id = input('Enter instructor id: ')
inst_name = input('Enter instructor name: ')
dept_name = input('Enter department name: ')
salary = input('Enter salary: ')

add_inst_query = """
 INSERT INTO instructor
 VALUES (%s, %s, %s, %s)
"""

val_tuple = (inst_id, inst_name, dept_name, salary)

with conn.cursor() as cursor:
 cursor.execute(add_inst_query, val_tuple)
 conn.commit()
```

```
Enter instructor id: 99999
Enter instructor name: Dave
Enter department name: Physics
Enter salary: 120000
```



# Inserting a Bulk of Records

- ▶ You can insert multiple records using the **executemany()** method
- ▶ It accepts two parameters:
  - ▶ A **query** that contains placeholders for the records that need to be inserted
  - ▶ A **list** that contains all records that you wish to insert

```
add_instructors_query = """
 INSERT INTO instructor
 (ID, name)
 VALUES (%s, %s)
"""

instructors = [
 ('00001', 'Mary'),
 ('00002', 'Thomas'),
 ('00003', 'Mitchell'),
 ('00004', 'Kimbra')
]

with conn.cursor() as cursor:
 cursor.executemany(add_instructors_query, instructors)
 conn.commit()
```

The screenshot shows the MySQL Workbench interface with a 'Result Grid' tab. The query 'SELECT \* FROM university.instructor;' has been run, and the results are displayed in a grid. The columns are labeled ID, name, dept\_name, and salary. The data includes rows for Mary, Thomas, Mitchell, and Kimbra, along with many other entries.

| ID    | name       | dept_name  | salary    |
|-------|------------|------------|-----------|
| 00001 | Mary       | NULL       | NULL      |
| 00002 | Thomas     | NULL       | NULL      |
| 00003 | Mitchell   | NULL       | NULL      |
| 00004 | Kimbra     | NULL       | NULL      |
| 10101 | Srinivasan | Comp. Sci. | 65000.00  |
| 12121 | Wu         | Finance    | 90000.00  |
| 15151 | Mozart     | Music      | 40000.00  |
| 22222 | Einstein   | Physics    | 95000.00  |
| 32343 | El Said    | History    | 60000.00  |
| 33456 | Gold       | Physics    | 87000.00  |
| 45565 | Katz       | Comp. Sci. | 75000.00  |
| 58583 | Calfieri   | History    | 62000.00  |
| 76543 | Singh      | Finance    | 80000.00  |
| 76766 | Crick      | Biology    | 72000.00  |
| 83821 | Brandt     | Comp. Sci. | 92000.00  |
| 98345 | Kim        | Elec. Eng. | 80000.00  |
| 99987 | Chris      | Comp. Sci. | 98000.00  |
| 99999 | Dave       | Physics    | 120000.00 |
| NULL  | NULL       | NULL       | NULL      |



# Update and Delete

- ▶ Updating and deleting work the same way, just pass the SQL to **cursor.execute()**
- ▶ You can use the cursor **rowcount** attribute to check how many records were affected:

```
delete_inst_query = """
 DELETE FROM instructor
 WHERE ID = '00001'
"""

with conn.cursor() as cursor:
 cursor.execute(delete_inst_query)
 conn.commit()
 print('Number of rows deleted:', cursor.rowcount)
```

Number of rows deleted: 1



# Calling Stored Procedures from Python

- ▶ To call a stored procedure use the **callproc()** method of the Cursor object  
`cursor.callproc(procedure_name, args=())`
- ▶ Then, you can call the **stored\_results()** method to get an iterator with the result set
  - ▶ The rows in the result can be read by calling the **fetchall()** method
- ▶ Example:

```
with conn.cursor() as cursor:
 cursor.callproc('get_courses_by_department', ('Comp. Sci.',))
 results = cursor.stored_results()
 for result in results:
 for course in result.fetchall():
 print(course)
```

```
('CS-101', 'Intro. to Computer Science', 'Comp. Sci.', Decimal('4'))
('CS-190', 'Game Design', 'Comp. Sci.', Decimal('4'))
('CS-315', 'Robotics', 'Comp. Sci.', Decimal('3'))
('CS-319', 'Image Processing', 'Comp. Sci.', Decimal('3'))
('CS-347', 'Database System Concepts', 'Comp. Sci.', Decimal('3'))
```



# Loading Data into a DataFrame

- ▶ Pandas provides a `read_sql()` method that reads an SQL query into a DataFrame

```
pandas.read_sql(sql, con, index_col=None, coerce_float=True, params=None, parse_dates=None,
columns=None, chunksize=None)
```

[source]

- ▶ For example, loading the instructor table into a DataFrame:

```
import pandas as pd

query = "SELECT * FROM instructor"
df = pd.read_sql(query, conn)
```

```
df.head()
```

|   | ID    | name       | dept_name  | salary  |
|---|-------|------------|------------|---------|
| 0 | 10101 | Srinivasan | Comp. Sci. | 65000.0 |
| 1 | 12121 | Wu         | Finance    | 90000.0 |
| 2 | 15151 | Mozart     | Music      | 40000.0 |
| 3 | 22222 | Einstein   | Physics    | 95000.0 |
| 4 | 32343 | El Said    | History    | 60000.0 |



# Loading Data into a DataFrame

- ▶ You can use the **params** argument to pass a list of parameters to the query
- ▶ For example, the following displays all the students that took the course CS-319

```
query = """
 SELECT * FROM takes
 WHERE course_id = %s
"""

df = pd.read_sql(query, conn, params=['CS-319'])
```

```
df
```

|   | ID    | course_id | sec_id | semester | year   | grade |
|---|-------|-----------|--------|----------|--------|-------|
| 0 | 45678 | CS-319    | 1      | Spring   | 2018.0 | B     |
| 1 | 76543 | CS-319    | 2      | Spring   | 2018.0 | A     |