

How many distinct last names of actors are there?

```
SELECT COUNT(DISTINCT last_name)
```

```
FROM actor;
```

Which actors participated in the movie 'Academy Dinosaur'? Print their first and last names.

```
SELECT first_name, last_name
```

```
FROM actor AS a
```

```
JOIN film_actor AS fa
```

```
ON a.actor_id = fa.actor_id
```

```
JOIN film AS f
```

```
ON f.film_id = fa.film_id
```

```
WHERE f.title = 'Academy Dinosaur';
```

3. How many copies of the film 'Hunchback Impossible' exist in the inventory system?

```
SELECT COUNT(*) AS copies_in_inventory
```

```
FROM inventory AS i
```

```
JOIN film AS f
```

```
ON i.film_id = f.film_id
```

```
WHERE f.title = 'Hunchback Impossible';
```

4. What is the total amount paid by each customer for all their rentals? For each customer print their name and the total amount paid.

```
SELECT c.first_name, c.last_name, SUM(p.amount) AS total_amount
```

```
FROM customer AS c
```

```
JOIN payment AS p
```

```
ON c.customer_id = p.customer_id
```

```
JOIN film AS f
```

```
ON f.film_id = p.film_id
```

```
WHERE f.title = 'Hunchback Impossible';
```

5. How many films from each category each store has? Print the store id, category name and number of films. Order the results by store id and category name.

```
SELECT i.store_id, c.name AS category, COUNT(*) AS films_num
```

```
FROM category AS c
```

```
JOIN film_category AS fc
```

```
ON c.category_id = fc.category_id
```

```
JOIN inventory AS i
```

```
ON i.film_id = fc.film_id
```

```
GROUP BY i.store_id, c.category_id
```

```
ORDER BY i.store_id, c.name;
```

6. Calculate the total revenue of each store.

```
SELECT i.store_id, SUM(amount) AS total_revenue
```

```
FROM payment AS p
```

```
JOIN rental AS r
```

```
ON r.rental_id = r.rental_id
```

```
JOIN inventory AS i
```

```
ON i.inventory_id = r.inventory_id
```

```
GROUP BY i.store_id;
```

7. Which actor participated in the most films? Print their full name and in how many movies they participated

```
SELECT first_name, last_name, COUNT(*) AS films_num
```

```
FROM actors AS a
```

```
JOIN film_actor AS fa
```

```
ON a.actor_id = fa.actor_id
```

```
GROUP BY a.actor_id
```

```
HAVING COUNT(*) >= ALL (
```

```
SELECT COUNT(*)
```

```
FROM film_actor
```

```
GROUP BY actor_id);
```

8. Find pairs of actors that participated together in the same movie and print their full names. Each such pair should appear only once in the result.

```
SELECT DISTINCT CONCAT(a1.first_name, ' ', a1.last_name) AS first_actor,
```

```
CONCAT(a2.first_name, ' ', a2.last_name) AS second_actor
```

```
FROM actor AS a1
```

```
JOIN film_actor AS fa1
```

```
ON a1.actor_id = fa1.actor_id
```

```
JOIN film_actor AS fa2
```

```
ON fa1.film_id = fa2.film_id
```

```
JOIN actor AS a2
```

```
ON a2.actor_id = fa2.actor_id
```

```
WHERE a1.actor_id < a2.actor_id;
```

9. Display the top five most popular films, i.e., films that were rented the highest number of times. For each film print its title and the number of times it was rented

```
SELECT title, COUNT(*) AS times_rented
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON f.film_id = i.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
GROUP BY f.film_id
```

```
ORDER BY times_rented DESC
```

```
LIMIT 5;
```

10. Is the film 'Academy Dinosaur' available for rent from Store 1? You should check that the film exists as one of the items in the inventory of Store 1, and that there is no outstanding rental of that item with no return date.

```
SELECT EXISTS (
```

```
SELECT *
```

```
FROM inventory AS i
```

```
JOIN film AS f
```

```
ON i.film_id = f.film_id
```

```
WHERE f.title = 'Academy Dinosaur'
```

```
AND i.store_id = 1
```

```
AND NOT EXISTS (
```

```
SELECT *
```

```
FROM rental AS r
```

```
WHERE i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
JOIN inventory AS i
```

```
ON i.film_id = f.film_id
```

```
JOIN rental AS r
```

```
ON i.inventory_id = r.inventory_id
```

```
AND r.return_date IS NULL) ) AS 'available';
```

```
FROM film AS f
```

```
Find how many flights departed from LAX airport in July 2015.
flights[flights['ORIGIN_AIRPORT'] == 'LAX'] & (flights['YEAR'] == 2015) &
(flights['MONTH'] == 7)
Find the number of the flight that had the longest arrival delay
flights.loc[flights['ARRIVAL_DELAY'].idxmax(), 'FLIGHT_NUMBER']
Find the airport with the highest number of arriving flights
flights.groupby('DESTINATION_AIRPORT').size().idxmax()
Find the day of week that had the highest number of flight cancellations.
flights.groupby('DAY_OF_WEEK')['CANCELLED'].sum().idxmax()
Create a bar plot showing the number of flights for each airline.
flights.groupby('AIRLINE').size().plot.bar()
Find the mean arrival delay for each airline.
flights.groupby('AIRLINE')['ARRIVAL_DELAY'].mean()
Find the airlines that had more than 10,000 cancellations.
df = flights[flights['CANCELLED'] == 1].groupby('AIRLINE').size()
df[df > 10000]
Find airlines having more than 2% of their flights cancelled. For each such airline,
print its identifier and the percentage of cancelled flights.
df = flights.groupby('AIRLINE')['CANCELLED'].mean()
df[df > 0.02]
Find the three top airlines with the highest number of cancelled or diverted flights.
flights[flights['CANCELLED'] == 1] |
(flights['DIVERTED'] == 1)][['AIRLINE'].value_counts()[:3]]
Find the longest sequence of on-time flights for each airline (an on-time flight is a
flight with less than 15 minutes arrival delay).
df['ontime'] = df['ARRIVAL_DELAY'] < 15
df.groupby('AIRLINE')['ontime'].apply(lambda x:
(~x).cumsum().where(x).value_counts().max())
Find the three airlines with the longest arrival delays.
Flights.groupby('Airline')['Arrival_Delay'].max().nlargest(3)
name and age of players who play for 'Boston Eagles' and their age is above 28.
df[(df['Team'] == 'Boston Eagles') & (df['Age'] > 28)][['Name', 'Age']]
the name and height of the tallest player.
df.loc[df['Height'].idxmax(), ['Name', 'Height']]
players whose salary is above the average salary in their team
df['Avg_Team_Salary'] = df.groupby(['Team'])['Salary'].transform('mean')
df['df_Salary'] >
df['Avg_Team_Salary']][['Name', 'Team', 'Salary', 'Avg_Team_Salary']]
```

```
Given an array A, print all the numbers that appear more than twice in the array.
values, counts = np.unique(A, return_counts=True)
values[counts > 2]
Given an array A and a number v, find the closest value in A to v
A[np.argmax(np.abs(A - v))].
Given a matrix A, print the number of rows in the matrix that contain the number 5
np.count_nonzero(np.any(A == 5, axis=1))
Given a matrix A, print the indices of the columns that contain a number greater
than 7
np.where(np.any(A > 7, axis=0))[0]
Given a matrix A, zero out all the rows that contain a negative number
A[np.where(np.any(A < 0, axis=1))] = np.zeros(A.shape[1])
Given a matrix A, normalize its rows such that the numbers in each row sum to 1.
A = A / np.sum(A, axis=1).reshape(-1, 1)
np.apply_along_axis(lambda row: row/np.sum(row) != 0 else row,
axis=1, arr = matrix)
Given an array A, find all the even numbers in the array that are followed by an odd
number.
B = A[: -1]
C = A[1:]
B[(B % 2 == 0) & (C % 2 == 1)]
Given an array A, find all the local maxima (peaks) in the array. A local maximum is
a number that is greater or equal to its two neighbors
B = A[1:-1] # the first and last numbers cannot be peaks
B[(B > A[: -2]) & (B > A[2:])]
Given a matrix A, find how many rows in A contain duplicate values.
B = np.sort(A, axis=1)
np.count_nonzero(np.any(B[:, 1:] == B[:, :-1], axis=1))
Given a matrix A, find all the saddle points in A
N = np.argmax(A, axis=1) # if change min -> max after
M = np.argmax(A, axis=0) # same
saddle_row = np.where(M[N] == np.arange(A.shape[0]))[0]
A[saddle_row, N[saddle_row]]
#latin square
def is_latin_square(matrix):
    n = matrix.shape[0]
    seq = np.arange(1, n + 1)
    rows_check = np.all(np.sort(matrix, axis=0) == seq.reshape(-1, 1))
    cols_check = np.all(np.sort(matrix, axis=1) == seq)
    return rows_check and cols_check
minimum = q1.min()
maximum = q1.max()
```

```
The average of every even row
answer_b = q1[:2].mean(axis=1)
The standard deviation in every odd column
answer_c = q1[:, 1::2].std(axis=0)
How many rows contain a number less than 5?
row_less5 = np.sum(np.any(q1<5, axis=1))
Which column contain a number greater than 90?
np.where(np.any(q1 > 90, axis=0))[0]
which numbers appear in the matrix more than twice?
values, counts = np.unique(q1, return_counts = True)
num_2 = values[counts > 2]
How many row contains duplicate value?
row_duplicate = np.sum(np.apply_along_axis(lambda x : len(np.unique(x)))=
len(x), axis = 1, arr = q1 )
#matplotlib
years = [1972, 1974, 1978, 1982, 1985, 1989, 1993, 1997, 1999, 2000, 2003,
2004, 2007, 2008, 2012]
transistors=[0.0024, 0.005, 0.029, 0.12, 0.275, 1.18, 3.1, 7.5, 24.0, 42.0, 220.0,
592.0, 1720.0, 2046.0, 3100.0]
y0 = 1972
n0 = transistors[0]
expected_transistors = [n0 * 2**( (year- y0)/2) for year in years]
log_actual = np.log(transistors)
log_expected = np.log(expected_transistors)
plt.figure(figsize = (8,4))
plt.plot(years, log_expected, '-.', label = "Moore's laws prediction")
plt.scatter(years, log_actual, marker = 'r', color='r', label = "Actual Transistor")
plt.xlabel("year")
plt.ylabel("Transistors(in millions) in log scale")
plt.legend()
plt.title("Moore's law vs actual transistors")
plt.grid(True)
```

```
browsers = ['Chrome', 'Safari', 'Edge', 'Firefox', 'Samsung Internet', 'Opera']
market_share = [65.18, 18.55, 5.26, 2.74, 2.56, 2.15]
colors = ['gold', 'silver', 'yellowgreen', 'lightcoral', 'lightskyblue', 'pink']
explode = (0, 1, 0, 0, 0, 0)
plt.title('Browser market shares worldwide(2024)')
plt.pie(market_share, explode=explode, labels=browsers, colors=colors, autopct
= '%1.1f%%', shadow=True, startangle=60, labeldistance=1.15)
plt.axis('equal')
plt.figure(figsize = (8,4))
plt.scatter(sepal_length, sepal_width, s = petal_length * 50, c = y, cmap = 'viridis',
edgecolor = 'k', alpha = 0.5)
plt.xlabel('sepal length(cm)')
plt.ylabel('sepal width(cm)')
plt.colorbar()
plt.grid(True)
plt.title("The Iris Flower Data Set")
```

```
fig, axes = plt.subplots(5, 10, figsize = (10, 5))
axes = axes.ravel()
for i in range(50):
axes[i].imshow(images[i], cmap='gray')
axes[i].set_title(label(i))
axes[i].axis('off')
plt.tight_layout()
from mpl_toolkits.mplot3d import Axes3D
x = np.linspace(-5, 5, 100)
y = np.linspace(-5, 5, 100)
X, Y = np.meshgrid(x, y)
Z = np.sin(X**2 + Y**2)
fig = plt.figure()
ax = plt.axes(projection = '3d')
surface = ax.plot_surface(X, Y, Z, cmap='coolwarm')
fig.colorbar(surface)
plt.show()
DBMS Architecture/ A DBMS consists of two major components
-Storage manager -Query processor
Good Database Design Principles/ Avoid duplicate information (redundancy)
Store information in its smallest logical parts Define a primary key for each table
If there is no column that might make a good primary key, use an auto-increment
column Ensure the integrity of the data by applying integrity constraints Apply
normalization rules
DCL(Data Definition Language) CREATE ALTER DROP RENAME TRUNCATE
DML(Data Manipulation Language) INSERT UPDATE DELETE MERGE CALL EXPLAIN
PLAN LOCK TABLE DQL(Data Query Language) SELECT DCL(Data Control Language)
GRANT REVOKE TCL(Transaction Control Language) COMMIT ROLLBACK
SAVEPOINT SET TRANSACTION
Primary Key
The primary key uniquely identifies each record in the table
[] A table can have only one primary key
[] The primary key can consist of one or more columns
[] The columns of the primary key are automatically defined as unique and not null
[] The primary key should be chosen such that its values are never/rarely changed
[] The primary key can be defined as part of a column declaration or appear on its
own
[] A foreign key is a column (or set of columns) in one table that refers to a primary
key
in another table
[] The table with the foreign key is called the child table
[] The table with the primary key is called the parent table
[] Foreign keys enforce referential integrity of the data
[] A value in the foreign key column must be one of the values contained in the
primary key
[] ON DELETE CASCADE - when a row in the parent table is deleted, automatically
delete its
corresponding child rows
[] ON UPDATE CASCADE - when the primary key column in the parent table is
updated,
automatically update the foreign key column in the corresponding child rows
[] ALTER TABLE allows you change the table definition a=ALTER TABLE table_name
[] To add a column to a table, use the following:\ a ADD column_name datatype;
[] To delete a column from a table: a DROP COLUMN column_name;
[] To change the data type of a column: a MODIFY COLUMN column_name datatype;
[] The DROP TABLE statement allows you to drop an existing table
[] The DROP table table_name; - Deletes both the data and the schema of the table
[] TRUNCATE TABLE deletes only the data in the table, but not the table itself
TRUNCATE TABLE table_name;
[] It's faster than DELETE if you want to delete all the records from the table
[] SQL DML commands allows you to manipulate the data in the database
[] Insertion of new rows into a given table
[] Deletion of rows in a given table[] Updating values in a given table
UPDATE employees
SET salary = salary * 1.1
WHERE employee_id = 1;
DELETE FROM employees
WHERE employee_id = 1;
Select
CASE
WHEN condition1 THEN result1
WHEN condition2 THEN result2
WHEN conditionN THEN resultN
ELSE
END (AS *)
The LIKE operator allows you to perform pattern matching on strings[] Can use two
wildcard characters for constructing patterns:[] percentage (%) matches any string
of 0 or more characters [] underscore (_) matches any single character
Select name from stu where name like 'B%'
```

```
axes[i].imshow(images[i], cmap='gray')
axes[i].set_title(label(i))
axes[i].axis('off')
plt.tight_layout()
```

```
from mpl_toolkits.mplot3d import Axes3D
x = np.linspace(-5, 5, 100)
y = np.linspace(-5, 5, 100)
X, Y = np.meshgrid(x, y)
Z = np.sin(X**2 + Y**2)
fig = plt.figure()
ax = plt.axes(projection = '3d')
surface = ax.plot_surface(X, Y, Z, cmap='coolwarm')
fig.colorbar(surface)
plt.show()
```

DBMS Architecture/ A DBMS consists of two major components

- Storage manager -Query processor

Good Database Design Principles/ Avoid duplicate information (redundancy)

Store information in its smallest logical parts Define a primary key for each table

If there is no column that might make a good primary key, use an auto-increment column

Ensure the integrity of the data by applying integrity constraints Apply normalization rules

DCL(Data Definition Language) CREATE ALTER DROP RENAME TRUNCATE

DML(Data Manipulation Language) INSERT UPDATE DELETE MERGE CALL EXPLAIN

PLAN LOCK TABLE DQL(Data Query Language) SELECT DCL(Data Control Language)

GRANT REVOKE TCL(Transaction Control Language) COMMIT ROLLBACK

SAVEPOINT SET TRANSACTION

Primary Key

The primary key uniquely identifies each record in the table

- [] A table can have only one primary key
- [] The primary key can consist of one or more columns
- [] The columns of the primary key are automatically defined as unique and not null
- [] The primary key should be chosen such that its values are never/rarely changed
- [] The primary key can be defined as part of a column declaration or appear on its own
- [] A foreign key is a column (or set of columns) in one table that refers to a primary key in another table
- [] The table with the foreign key is called the child table
- [] The table with the primary key is called the parent table
- [] Foreign keys enforce referential integrity of the data
- [] A value in the foreign key column must be one of the values contained in the primary key
- [] ON DELETE CASCADE - when a row in the parent table is deleted, automatically delete its corresponding child rows
- [] ON UPDATE CASCADE - when the primary key column in the parent table is updated, automatically update the foreign key column in the corresponding child rows
- [] ALTER TABLE allows you change the table definition a=ALTER TABLE table\_name
- [] To add a column to a table, use the following:\ a ADD column\_name datatype;
- [] To delete a column from a table: a DROP COLUMN column\_name;
- [] To change the data type of a column: a MODIFY COLUMN column\_name datatype;
- [] The DROP TABLE statement allows you to drop an existing table
- [] The DROP table table\_name; - Deletes both the data and the schema of the table
- [] TRUNCATE TABLE deletes only the data in the table, but not the table itself TRUNCATE TABLE table\_name;
- [] It's faster than DELETE if you want to delete all the records from the table
- [] SQL DML commands allows you to manipulate the data in the database
- [] Insertion of new rows into a given table
- [] Deletion of rows in a given table[] Updating values in a given table
- UPDATE employees
- SET salary = salary \* 1.1
- WHERE employee\_id = 1;

```
DELETE FROM employees
WHERE employee_id = 1;
```

```
Select
CASE
WHEN condition1 THEN result1
WHEN condition2 THEN result2
WHEN conditionN THEN resultN
ELSE
END (AS *)
```

The LIKE operator allows you to perform pattern matching on strings[] Can use two wildcard characters for constructing patterns:[] percentage (%) matches any string of 0 or more characters [] underscore (\_) matches any single character

Select name from stu where name like 'B%'

NULL Values

- [] A field with a NULL value is a field that has no value
- [] The result of any arithmetic expression involving NULL is NULL
- [] Comparing NULL values with operators such as =, < returns an unknown result
- [] You have to use the IS NULL and IS NOT NULL operators instead

- [] You can combine the result sets of two or more queries using set operators
- [] Every select must have the same number of columns
- [] The corresponding columns must have similar data types

SELECT column\_name(s) FROM table1

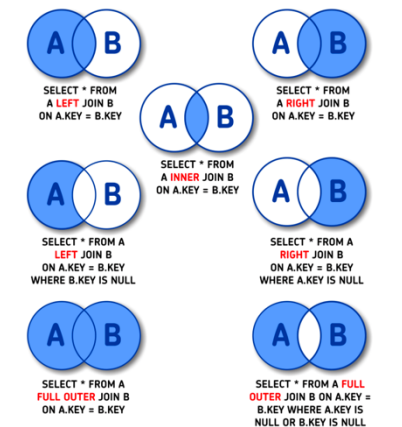
UNION

SELECT column\_name(s) FROM table2;

All aggregate functions except COUNT(\*) ignore NULL values in their input

[] COUNT(\*) returns the number of rows in the table, regardless of NULL values

[] COUNT(column) returns the number of non-null values in the specified column



SQL EXECUTION ORDER: From, on, join, where, group by, with, having, select, distinct, order by, top

CTE is a named temporary result set that only exists for the duration of the query

```
WITH courses_per_student AS (
-- Count the number of courses each student took
SELECT id, COUNT(DISTINCT course_id) AS num_courses
```

```
FROM table
GROUP BY id
)
-- Get the distribution of courses taken by students
SELECT num_courses, COUNT(id) AS num_students
FROM courses_per_student
GROUP BY num_courses
ORDER BY num_courses;
```

A view is a virtual table based on the results set of an SQL statement

- [] Reasons to use views:
- [] Data security - allows you to hide certain data from view of certain users
- [] e.g., some users only need to know the instructor IDs and names, but not their salaries
- [] Hiding complexity of the underlying tables
- [] A view is created with the CREATE VIEW statement: CREATE VIEW view\_name AS SELECT column1, column2, ... FROM table\_name WHERE condition;
- [] A view always shows up-to-date data
- [] The database engine recreates the view whenever it is used in a query

- [] In order to execute SQL queries in Python, you need a cursor object. A cursor object allows you to traverse over database records. To process the rows in the result one at a time, you can use the cursor as an iterator: The cursor has a few useful properties that provide information about the result set
- [] column\_names - returns the list containing the column names of the result set
- [] rowcount - the number of rows in the result set

A parameterized query is a query that uses placeholders (%) for attribute values(solution to injection attack)

A transaction is a sequence of SQL statements that represents a single unit of work(The transaction must end with one of the following statements:

- [] COMMIT: The updates performed by the transaction become permanent in the database
- [] ROLLBACK: All updates performed by the transaction are undone transactions need to be written inside a stored procedure

DECLARE EXIT HANDLER FOR SQLWARNING ROLLBACK;

START TRANSACTION; COMMIT;

An index is a data structure that can speed up queries/ searches Causes updates to the table to become slower (since the index also needs to be updated), Storage Space

CREATE [UNIQUE] INDEX index\_name ON table\_name (column1, column2, ...); A clustered index defines the physical order in which table records are stored There can be only one clustered index per table. By default a clustered index is created on a primary key column

Numpy: To create a 2D array, you can pass a list of lists. A=np.array([1,2],[3,4]) A=np.zeros((3,3)), B=np.ones(5),C=np.full((3,3),7).

np.random.random(shape=(3,3)), creates an array with numbers sampled randomly from the uniform distribution over [0, 1). np.random.randint(low, [high], size=(1,10,(5,5)) creates an array with integer numbers sampled from the interval [low, high) np.random.normal(loc=0.0, scale=1.0, size=(3,3)) samples numbers from the normal distribution with mean = loc and standard deviation = scale. np.arange(start, stop, [step]) creates an array from a sequence of numbers. Its arguments are the same as in Python's range() function. np.arange(0,2,0.2)->array([0., 0.2, 0.4, 0.6, 0.8, 1., 1.2, 1.4, 1.6, 1.8]) np.linspace(start, stop, num) creates an array of evenly spaced numbers np.linspace(0.5,4)->array([0., 1.66666667, 3.33333333, 5.]) reshape(3,4)-> will change the dimension to 3 rows and 4 columns. ravel() can be used to flatten a multidimensional array into one dimension: array.ravel()

Indexing: A[1:2] # row 1, column 2. A[-1,-1] # last row and column. A[1] # all row 1. Slicing A[0:2] # row 0 to 2 not included (first two). A[:0:2] # column 0 to 2 not included (first two). A[0:2:0:2] # the first two rows and column. a[1,1:]# the second row and second column onwards. Fancy indexing: a[[0,2]] first and third row. A[:,1:3]] second and fourth columns. Boolean indexing/masking is used to select elements of the array that satisfy some condition: a[a%2==0] will return the numbers in a. a%2==0 will return true/false masking. Changing elements: Changes to a sliced or indexed array are reflected in the original array. a[[2,4]]-=99 or a[a<5]->99-> If a number is equal to 2 or 4 it will change to 99. Comparisons are performed element-wise and produce an array of Boolean values. Logical operators: & | Calculations by row and column a.mean(axis=0). Missing values: Do not test nans for equality (np.nan == np.nan is False), instead use np.isnan(a)->result will be True/False. NaN is a bit like a virus - the result of any arithmetic with nan will also be a nan: if you have nan the result for the calculation would be nan. Eg. np.sum(), np.nansum() is a safe NaN version. np.any(a>5,axis=0) or np.all(). np.sort(a) - will not sort the original. a.sort() - will. Can add axis. np.argsort() returns the indexes that would sort an array rather than the sorted elements themselves

Searching: np.where(condition) returns the indices of all the elements that satisfy the condition.In case of a 2D array, it returns a tuple of the row and the column indices of the elements that satisfy the condition. Unique values: np.unique()-> value, count = np.unique(a, return\_counts=True) + value[count>2] -> result will be the number which appear more than 2. Closest value in the array: def closest(a,v): /dis = np.abs(a-v) / return a[np.argmin(dis)]. np.concatenate([100],array,[100]) np.append(arr, values, axis=None)

appends values to the end of an array. values must be of the correct shape (the same shape as arr, excluding axis). axis specifies the axis along which to insert values. np.append(a,[1,2,3], axis = 0). np.vstack((a,b)) stacks the arrays vertically (row-wise). np.hstack((a,b)) stacks the arrays horizontally (column-wise). Broadcasting allows operations to be performed on arrays of different shapes. Rule 1: If the two arrays have different numbers of dimensions, the shape of the smaller array is padded with ones on its left side. Rule 2: The arrays are compatible if all their dimensions are equal or one of them is 1. Rule 3: If the arrays are compatible and their shapes don't match in one of the dimensions, the array with dimension of 1 is stretched to match the other array.

np.eye(N) creates an identity matrix of size N xN. np.diag(v) creates a diagonal matrix from the 1-D array v: The transpose of a matrix results from "flipping" the rows and columns. Did not include: Linear Algebra, Linear Equations, Determinant, Eigenvectors and Eigenvalues) Pandas: Can also construct a series from a dictionary- just insert the name of the dictionary. df.index gives the row labels.df.columns give the column labels.df.dtypes gives the data types of the columns df.shape gives the axis dimensions as in Numpy. values gives the data in the DataFrame as a 2D Numpy array.We can access rows by its index or name: ser[2] or ser[0]. Slicing: same as accessing-> ser[0:2], ser['a':'b'] only the first two. df.loc[] selects a group of rows and columns using the index and column names. A single label, e.g., 'a' A list of labels, e.g., ['a', 'b'] A slice object with labels, e.g., 'a':'f' A Boolean mask, e.g., 'a' > 5. df.iloc[] selects a group of rows and columns based on their integer positions. An integer, e.g., 5 A list of integers, e.g., [4, 3, 0] A slice object with ints, e.g., 1:7 You can use loc and iloc to assign values to specific cells: df.loc['age'] += 1. To select a single column you can also use df[column\_name] or df[column\_name]. d = df.loc[(df['animal'] == 'cat') & (df['size'] == 'L')] ... (df.loc[df.index[-1], 'weight'] = 15) or df.iloc[-1, df.columns.get\_loc('weight')] = 15) You can add a column to the DataFrame by assigning a list or a Numpy array to a new column name: df.new\_column\_name = [...] sort\_index(axis=0, ascending=True) sorts the DataFrame by its row or column index. Use axis=0 to sort by the row index and axis=1 to sort by the column names. sort\_values(by=axis=0, ascending=True) sorts by the values along the specified axis by a label or list of labels to sort by. nlargest(n, columns) returns the n rows with the largest values in the specified columns. isna(), isnull() = Generate a boolean mask indicating missing values. dropna(axis=0, how='any') = Return a filtered version of the data. fillna(S) = Return a copy of the data with missing values filled or imputed.